

Linear models

Modeling... linearly!

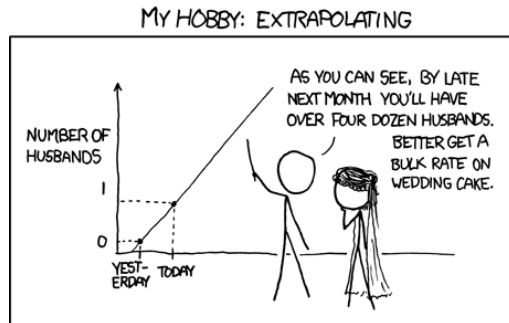
Samuel Robinson, Ph.D.

Sep. 22, 2023

Part 1: How do they work?

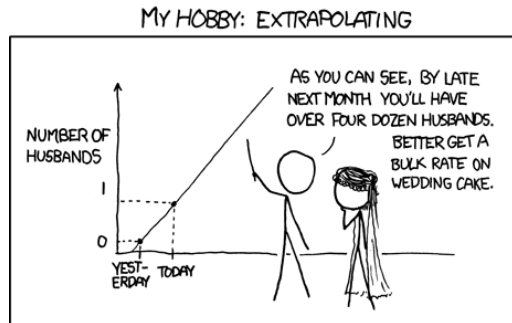
Outline

- What are linear models? How do I fit them?



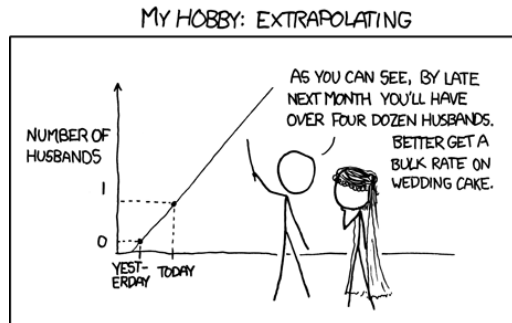
Outline

- What are linear models? How do I fit them?
- Making sure the model is working properly



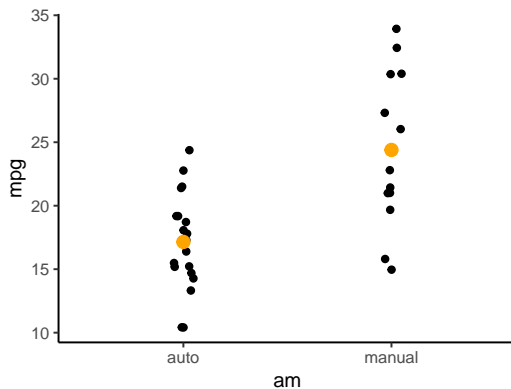

Outline

- What are linear models? How do I fit them?
- Making sure the model is working properly
- Plotting and interpreting model results



Motivation

- *I measured 2 things and I want to know if they're related to each other*
- *I have groups of data, and I want to know whether the means are different*



Terminology

Linear models go by many different names. All these models are all doing *exactly the same thing*:

- Linear regression

Terminology

Linear models go by many different names. All these models are all doing *exactly the same thing*:

- Linear regression
- Least-squares regression

Terminology

Linear models go by many different names. All these models are all doing *exactly the same thing*:

- Linear regression
- Least-squares regression
- Simple linear model (SLM)

Terminology

Linear models go by many different names. All these models are all doing *exactly the same thing*:

- Linear regression
- Least-squares regression
- Simple linear model (SLM)
- Multiple linear model/regression

Terminology

Linear models go by many different names. All these models are all doing *exactly the same thing*:

- Linear regression
- Least-squares regression
- Simple linear model (SLM)
- Multiple linear model/regression
- Analysis of Variance (ANOVA)

Terminology

Linear models go by many different names. All these models are all doing *exactly the same thing*:

- Linear regression
- Least-squares regression
- Simple linear model (SLM)
- Multiple linear model/regression
- Analysis of Variance (ANOVA)
- Analysis of Covariance (ANCOVA)

Terminology

Linear models go by many different names. All these models are all doing *exactly the same thing*:

- Linear regression
- Least-squares regression
- Simple linear model (SLM)
- Multiple linear model/regression
- Analysis of Variance (ANOVA)
- Analysis of Covariance (ANCOVA)

Terminology

Linear models go by many different names. All these models are all doing *exactly the same thing*:

- Linear regression
- Least-squares regression
- Simple linear model (SLM)
- Multiple linear model/regression
- Analysis of Variance (ANOVA)
- Analysis of Covariance (ANCOVA)

I use a set of terminology that I find very helpful, from [Berliner \(1996\)](#). I'll be using it here, as well as for describing more complex models.

Model terminology

All linear models take the form:

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 \dots + b_ix_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

Model terminology

All linear models take the form:

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 \dots + b_ix_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is the thing you're interested in predicting

Model terminology

All linear models take the form:

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 \dots + b_ix_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is the thing you're interested in predicting
- \hat{y} is the *predicted value* of y

Model terminology

All linear models take the form:

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 \dots + b_ix_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is the thing you're interested in predicting
- \hat{y} is the *predicted value* of y
- $x_1 \dots x_i$ are *predictors* of y

Model terminology

All linear models take the form:

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 \dots + b_jx_j$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is the thing you're interested in predicting
- \hat{y} is the *predicted value* of y
- $x_1 \dots x_j$ are *predictors* of y
- $b_1 \dots b_j$ are *coefficients* for each predictor x_j

Model terminology

All linear models take the form:

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 \dots + b_ix_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is the thing you're interested in predicting
- \hat{y} is the *predicted value* of y
- $x_1 \dots x_i$ are *predictors* of y
- $b_1 \dots b_i$ are *coefficients* for each predictor x_i
- b_0 is the *intercept*, a coefficient that doesn't depend on predictors

Model terminology

All linear models take the form:

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 \dots + b_ix_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is the thing you're interested in predicting
- \hat{y} is the *predicted value* of y
- $x_1 \dots x_i$ are *predictors* of y
- $b_1 \dots b_i$ are *coefficients* for each predictor x_i
- b_0 is the *intercept*, a coefficient that doesn't depend on predictors
- $y \sim \text{Normal}(\hat{y}, \sigma)$ means:

Model terminology

All linear models take the form:

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 \dots + b_ix_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is the thing you're interested in predicting
- \hat{y} is the *predicted value* of y
- $x_1 \dots x_i$ are *predictors* of y
- $b_1 \dots b_i$ are *coefficients* for each predictor x_i
- b_0 is the *intercept*, a coefficient that doesn't depend on predictors
- $y \sim \text{Normal}(\hat{y}, \sigma)$ means:
 - “ y follows a Normal distribution with mean \hat{y} and SD σ ”

Model terminology

All linear models take the form:

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 \dots + b_ix_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is the thing you're interested in predicting
- \hat{y} is the *predicted value* of y
- $x_1 \dots x_i$ are *predictors* of y
- $b_1 \dots b_i$ are *coefficients* for each predictor x_i
- b_0 is the *intercept*, a coefficient that doesn't depend on predictors
- $y \sim \text{Normal}(\hat{y}, \sigma)$ means:
 - “ y follows a Normal distribution with mean \hat{y} and SD σ ”

Model terminology

All linear models take the form:

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 \dots + b_ix_i$$


$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is the thing you're interested in predicting
- \hat{y} is the *predicted value* of y
- $x_1 \dots x_i$ are *predictors* of y
- $b_1 \dots b_i$ are *coefficients* for each predictor x_i
- b_0 is the *intercept*, a coefficient that doesn't depend on predictors
- $y \sim \text{Normal}(\hat{y}, \sigma)$ means:
 - “ y follows a Normal distribution with mean \hat{y} and SD σ ”

This may look terrifying, but let's use a simple example:


Example

- *mpg* is the thing you're interested in predicting



$$\hat{mpg} = b_0 + b_1 disp$$


Example



- *mpg* is the thing you're interested in predicting
- \hat{mpg} is the *predicted value* of *mpg*

$$\hat{mpg} = b_0 + b_1 disp$$


Example



- *mpg* is the thing you're interested in predicting
- \hat{mpg} is the *predicted value* of *mpg*
- *disp* is the *predictor* of *mpg*

$$\hat{mpg} = b_0 + b_1 disp$$


Example



- mpg is the thing you're interested in predicting
- \hat{mpg} is the *predicted value* of mpg
- $disp$ is the *predictor* of mpg
- b_0 is the *intercept*, b_1 is the *coefficient* (slope) for $disp$

$$\hat{mpg} = b_0 + b_1 disp$$


Example



- mpg is the thing you're interested in predicting
- \hat{mpg} is the *predicted value* of mpg
- $disp$ is the *predictor* of mpg
- b_0 is the *intercept*, b_1 is the *coefficient* (slope) for $disp$
- $mpg \sim Normal(\hat{mpg}, \sigma)$ means:

$$\hat{mpg} = b_0 + b_1 disp$$


Example



$$\hat{mpg} = b_0 + b_1 disp$$

- mpg is the thing you're interested in predicting
- \hat{mpg} is the *predicted value* of mpg
- $disp$ is the *predictor* of mpg
- b_0 is the *intercept*, b_1 is the *coefficient* (slope) for $disp$
- $mpg \sim Normal(\hat{mpg}, \sigma)$ means:
 - “ mpg follows a Normal distribution with mean \hat{mpg} and SD σ ”

Example




$$\hat{mpg} = b_0 + b_1 disp$$

- *mpg* is the thing you're interested in predicting
- \hat{mpg} is the *predicted value* of *mpg*
- *disp* is the *predictor* of *mpg*
- b_0 is the *intercept*, b_1 is the *coefficient* (slope) for *disp*
- $mpg \sim Normal(\hat{mpg}, \sigma)$ means:
 - “*mpg* follows a Normal distribution with mean \hat{mpg} and SD σ ”
- σ isn't displayed on the figure. Where is it?

Example (cont.)


σ isn't displayed on the figure. Where is it?

- σ is the “leftover” or “residual” variance



Example (cont.)


σ isn't displayed on the figure. Where is it?



- σ is the “leftover” or “residual” variance
- i.e. variation between samples that the model couldn't explain

Example (cont.)


σ isn't displayed on the figure. Where is it?



- σ is the “leftover” or “residual” variance
- i.e. variation between samples that the model couldn't explain
- Since $y \sim \text{Normal}(\hat{y}, \sigma)$, this means that points are normally distributed around the *entire line* of \hat{y}

Example (cont.)

σ isn't displayed on the figure. Where is it?



- σ is the “leftover” or “residual” variance
- i.e. variation between samples that the model couldn't explain
- Since $y \sim \text{Normal}(\hat{y}, \sigma)$, this means that points are normally distributed around the *entire line* of \hat{y}
- If you took a vertical slice at each part of the x-axis, the distribution would be *Normal*

How do I get R to fit this model?

lm is one of the main functions used for linear modeling:

```
#Formula= y ~ x, data = Name of the dataframe containing mpg & disp  
mod1 <- lm(mpg ~ disp, data = mtcars); summary(mod1)
```

```
##  
## Call:  
## lm(formula = mpg ~ disp, data = mtcars)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -4.8922 -2.2022 -0.9631  1.6272  7.2305   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept) 29.599855   1.229720  24.070  < 2e-16 ***  
## disp        -0.041215   0.004712  -8.747  9.38e-10 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 3.251 on 30 degrees of freedom  
## Multiple R-squared:  0.7183, Adjusted R-squared:  0.709   
## F-statistic: 76.51 on 1 and 30 DF,  p-value: 9.38e-10
```

For a detailed breakdown of lm's output, click [here](#)

Simulate data

Now that we know how linear models work, we can simulate our own data:

```
#Parameters:
```

```
b0 <- 1 #Intercept
```

```
b1 <- 2 #Slope
```

```
sigma <- 3 #SD
```

```
#Make up some data:
```

```
x <- 0:30 #Predictor values
```

```
#Predicted y values
```


```
pred_y <- b0 + b1*x
```

```
#Add "noise" around pred_y
```

```
actual_y <- rnorm(n = length(pred_y),  
                 mean = pred_y,  
                 sd= sigma)
```

```
#Plot the data we just made
```

```
plot(x,pred_y,col='orange',pch=19,type='l',  
      ylab='y')  
points(x,actual_y,col='black',pch=19)
```



Fit a model from simulated data

How does R do at finding the coefficients? Remember: $b_0 = 1$, $b_1 = 2$, $\sigma = 3$

```
fakeDat <- data.frame(x = x, y = actual_y, pred = pred_y) #Simulated data in a dataframe  
mod1sim <- lm(y ~ x, data = fakeDat); summary(mod1sim) #Fit model
```

```
##  
## Call:  
## lm(formula = y ~ x, data = fakeDat)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -5.7568 -1.7623 -0.2176  1.9419  5.3572   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  2.02974     1.00445   2.021  0.0526 .      
## x            1.92670     0.05751  33.499 <2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 2.864 on 29 degrees of freedom  
## Multiple R-squared:  0.9748, Adjusted R-squared:  0.9739   
## F-statistic: 1122 on 1 and 29 DF,  p-value: < 2.2e-16
```

Fit a model from simulated data

How does R do at finding the coefficients? Remember: $b_0 = 1$, $b_1 = 2$, $\sigma = 3$


```
fakeDat <- data.frame(x = x, y = actual_y, pred = pred_y) #Simulated data in a dataframe  
mod1sim <- lm(y ~ x, data = fakeDat); summary(mod1sim) #Fit model
```

```
##  
## Call:  
## lm(formula = y ~ x, data = fakeDat)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -5.7568 -1.7623 -0.2176  1.9419  5.3572   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  2.02974     1.00445   2.021  0.0526 .      
## x            1.92670     0.05751  33.499 <2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 2.864 on 29 degrees of freedom  
## Multiple R-squared:  0.9748, Adjusted R-squared:  0.9739   
## F-statistic: 1122 on 1 and 29 DF,  p-value: < 2.2e-16
```

Modeling philosophy: all models are approximating a **generative process**.

It is up to us to think about what this process might be like.

What about categorical data?




This uses *exactly the same* math!

- *mpg* is the thing you're interested in predicting

$$\hat{mpg} = b_0 + b_1 am$$

What about categorical data?




This uses *exactly the same math!*

- *mpg* is the thing you're interested in predicting
- \hat{mpg} is the *predicted value* of *mpg*

$$\hat{mpg} = b_0 + b_1 am$$

What about categorical data?




This uses *exactly the same math!*

- *mpg* is the thing you're interested in predicting
- \hat{mpg} is the *predicted value* of *mpg*
- *am* is the *predictor* of *mpg*

$$\hat{mpg} = b_0 + b_1 am$$

What about categorical data?




This uses *exactly the same* math!

- mpg is the thing you're interested in predicting
- \hat{mpg} is the *predicted value* of mpg
- am is the *predictor* of mpg
 - set of 0s and 1s, not continuous

$$\hat{mpg} = b_0 + b_1 am$$

What about categorical data?




This uses *exactly the same math!*

- mpg is the thing you're interested in predicting
- \hat{mpg} is the *predicted value* of mpg
- am is the *predictor* of mpg
 - set of 0s and 1s, not continuous
- b_0 is the *intercept*, b_1 is the *coefficient* for am

$$\hat{mpg} = b_0 + b_1 am$$

What about categorical data?



$$\hat{mpg} = b_0 + b_1 am$$

This uses *exactly the same* math!

- mpg is the thing you're interested in predicting
- \hat{mpg} is the *predicted value* of mpg
- am is the *predictor* of mpg
 - set of 0s and 1s, not continuous
- b_0 is the *intercept*, b_1 is the *coefficient* for am
- Where is σ ?

How do I get R to fit this model?

Syntax is exactly the same for this model

```
#Formula structure: y ~ x  
mod2 <- lm(mpg ~ am, #mpg depends on am  
           data = mtcars) #Name of the dataframe containing mpg & am  
summary(mod2)
```

```
##  
## Call:  
## lm(formula = mpg ~ am, data = mtcars)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -9.3923 -3.0923 -0.2974  3.2439  9.5077   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)   17.147      1.125   15.247 1.13e-15 ***  
## am             7.245      1.764    4.106 0.000285 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 4.902 on 30 degrees of freedom  
## Multiple R-squared:  0.3598, Adjusted R-squared:  0.3385   
## F-statistic: 16.86 on 1 and 30 DF,  p-value: 0.000285
```

A challenger approaches!

- Simulate your own data with 2 discrete levels. My suggestion:

A challenger approaches!

- Simulate your own data with 2 discrete levels. My suggestion:
 - ~~Steal~~Borrow my code, and change the predictor from continuous to discrete

A challenger approaches!

- Simulate your own data with 2 discrete levels. My suggestion:
 - Steal/Borrow my code, and change the predictor from continuous to discrete
 - Useful command: `rep` (replicate)

A challenger approaches!

- Simulate your own data with 2 discrete levels. My suggestion:
 - Steal/Borrow my code, and change the predictor from continuous to discrete
 - Useful command: `rep` (replicate)
 - e.g. `rep(x=c(0,1),each=10)`

A challenger approaches!

- Simulate your own data with 2 discrete levels. My suggestion:
 - Steal/Borrow my code, and change the predictor from continuous to discrete
 - Useful command: `rep` (replicate)
 - e.g. `rep(x=c(0,1),each=10)`
 - Useful command: `rnorm` (generate normally-distributed data)

A challenger approaches!

- Simulate your own data with 2 discrete levels. My suggestion:
 - ~~Steal~~Borrow my code, and change the predictor from continuous to discrete
 - Useful command: `rep` (replicate)
 - e.g. `rep(x=c(0,1),each=10)`
 - Useful command: `rnorm` (generate normally-distributed data)
 - e.g. `rnorm(n=100,mean=0,sd=1)`

A challenger approaches!

- Simulate your own data with 2 discrete levels. My suggestion:
 - ~~Steal~~Borrow my code, and change the predictor from continuous to discrete
 - Useful command: `rep` (replicate)
 - e.g. `rep(x=c(0,1),each=10)`
 - Useful command: `rnorm` (generate normally-distributed data)
 - e.g. `rnorm(n=100,mean=0,sd=1)`
- Use `lm` to fit a model to the data you just simulated



A challenger approaches!

- Simulate your own data with 2 discrete levels. My suggestion:
 - ~~Steal~~Borrow my code, and change the predictor from continuous to discrete
 - Useful command: `rep` (replicate)
 - e.g. `rep(x=c(0,1),each=10)`
 - Useful command: `rnorm` (generate normally-distributed data)
 - e.g. `rnorm(n=100,mean=0,sd=1)`
- Use `lm` to fit a model to the data you just simulated
 - How does R do at guessing your coefficients?


Part 2: More bells and whistles

Motivation

- *I have 2+ groups of data, and I want to know whether the means are different*
- *I have 2+ groups of bivariate data, and I want to know whether the relationships differ between groups*



Categorical data, 3 categories



The more factor levels, the more coefficients:

- mpg is the thing you're interested in predicting
- \hat{mpg} is the *predicted value* of mpg
- $gear$ is the *predictor* of mpg
- set of 0s and 1s
- $gears_4$ = "is this data point from a 4-gear car?"
- b_0 = *intercept*
- $[b_1, b_2]$ = are *coefficients* for $gears$

$$\hat{mpg} = b_0 + b_1 gears_4 + b_2 gears_5$$

How do I get R to fit this model?

```
#Formula structure: y ~ x  
mod1 <- lm(mpg ~ factor(gear), #mpg depends on gears  
           data = mtcars) #Name of the dataframe containing mpg & gears  
summary(mod1)
```


```
##  
## Call:  
## lm(formula = mpg ~ factor(gear), data = mtcars)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -6.7333 -3.2333 -0.9067  2.8483  9.3667   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)    16.107      1.216   13.250 7.87e-14 ***  
## factor(gear)4     8.427      1.823    4.621 7.26e-05 ***  
## factor(gear)5     5.273      2.431    2.169  0.0384 *    
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 4.708 on 29 degrees of freedom  
## Multiple R-squared:  0.4292, Adjusted R-squared:  0.3898   
## F-statistic: 10.9 on 2 and 29 DF,  p-value: 0.0002948
```

Dummy variables

```
mod1Matrix <- model.matrix(mod1) #Get model matrix (columns used to predict mpg)
head(mod1Matrix,28) #Show first 28 rows of model matrix
```

```
##                (Intercept) factor(gear)4 factor(gear)5
## Mazda RX4                1              1              0
## Mazda RX4 Wag            1              1              0
## Datsun 710                 1              1              0
## Hornet 4 Drive            1              0              0
## Hornet Sportabout         1              0              0
## Valiant                   1              0              0
## Duster 360                1              0              0
## Merc 240D                 1              1              0
## Merc 230                   1              1              0
## Merc 280                   1              1              0
## Merc 280C                  1              1              0
## Merc 450SE                 1              0              0
## Merc 450SL                 1              0              0
## Merc 450SLC                1              0              0
## Cadillac Fleetwood        1              0              0
## Lincoln Continental        1              0              0
## Chrysler Imperial         1              0              0
## Fiat 128                   1              1              0
## Honda Civic                1              1              0
## Toyota Corolla             1              1              0
## Toyota Corona              1              0              0
## Dodge Challenger           1              0              0
## AMC Javelin                 1              0              0
## Camaro Z28                  1              0              0
## Pontiac Firebird           1              0              0
## Fiat X1-9                   1              1              0
## Porsche 914-2              1              0              1
## Lotus Europa                1              0              1
```

What about if 2 things are both important?



- Suppose that both *disp* and *gears* are important for predicting *mpg*?
- This is very similar to the last example, except that now we've added *disp*
- *gears* now changes the intercept, while *disp* changes the slope of all the lines
- Does it look like *gear* is very important?

$$\hat{mpg} = b_0 + b_1 disp$$

How do I get R to fit this model?

#mpg depends on disp and gears

```
mod2 <- lm(mpg ~ disp+factor(gear), data = mtcars)
summary(mod2)
```

```
##
## Call:
## lm(formula = mpg ~ disp + factor(gear), data = mtcars)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
##	-4.9155	-2.1892	-0.9054	1.5790	7.2498

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	29.411183	2.627966	11.192	7.58e-12 ***
## disp	-0.040774	0.007601	-5.364	1.03e-05 ***
## factor(gear)4	0.138017	2.021332	0.068	0.946
## factor(gear)5	0.224712	1.976090	0.114	0.910

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.365 on 28 degrees of freedom
## Multiple R-squared:  0.7185, Adjusted R-squared:  0.6883
## F-statistic: 23.82 on 3 and 28 DF,  p-value: 7.31e-08
```

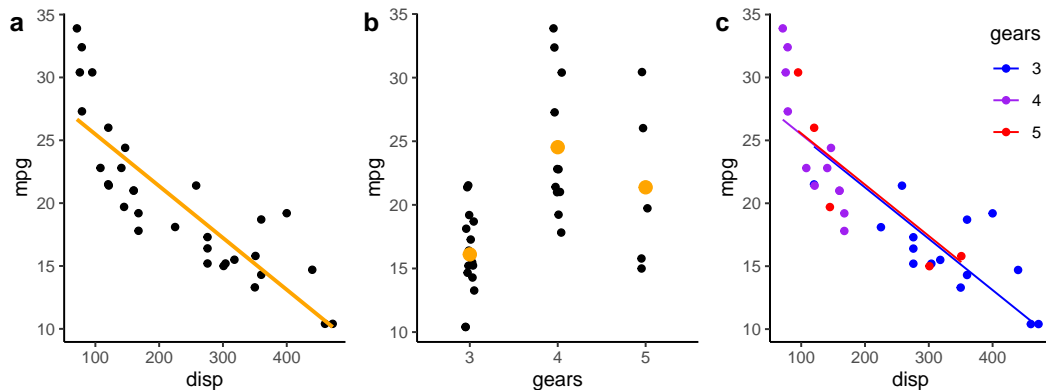
Dummy variables

```
mod2Matrix <- model.matrix(mod2) #Get model matrix (columns used to predict mpg)
colnames(mod2Matrix) <- gsub('factor\\(gear\\)', 'gear', colnames(mod2Matrix)) #Shorten colnames
head(mod2Matrix, 28) #Show first 28 rows of model matrix
```

```
##              (Intercept)  disp gear4 gear5
## Mazda RX4              1 160.0      1      0
## Mazda RX4 Wag          1 160.0      1      0
## Datsun 710              1 108.0      1      0
## Hornet 4 Drive          1 258.0      0      0
## Hornet Sportabout       1 360.0      0      0
## Valiant                 1 225.0      0      0
## Duster 360              1 360.0      0      0
## Merc 240D               1 146.7      1      0
## Merc 230                1 140.8      1      0
## Merc 280                1 167.6      1      0
## Merc 280C               1 167.6      1      0
## Merc 450SE              1 275.8      0      0
## Merc 450SL              1 275.8      0      0
## Merc 450SLC             1 275.8      0      0
## Cadillac Fleetwood      1 472.0      0      0
## Lincoln Continental     1 460.0      0      0
## Chrysler Imperial       1 440.0      0      0
## Fiat 128                 1  78.7      1      0
## Honda Civic              1  75.7      1      0
## Toyota Corolla          1  71.1      1      0
## Toyota Corona           1 120.1      0      0
## Dodge Challenger        1 318.0      0      0
## AMC Javelin             1 304.0      0      0
## Camaro Z28              1 350.0      0      0
## Pontiac Firebird         1 400.0      0      0
## Fiat X1-9                1  79.0      1      0
## Porsche 914-2           1 120.3      0      1
```

Interlude: problems with plotting raw data

- Say that I've fit the following model:
 $\text{mpg} \sim \text{disp} + \text{factor}(\text{gear})$
- All of the plots below are using raw data, but which one is “telling the truth”?
- Answer: **c**. *a* and *b* are hiding the effect of the other variable



How do I plot these model results?


Rule for plotting model results:

- 1 If the model uses N variables, you should show all N effects *simultaneously*

Other names for partial effects:

Incorrect example, using raw data:

```
#Fit model with 5 variables (all important)
simMod <- lm(y~X1+X2+X3+X4+X5,data=pred)
#Incorrect way, using raw data
plot(y~X5,data=pred,pch=19,cex.lab=3)
```



The effect of $X5$ is actually **very** strong ($p < 0.0001$), but it doesn't look like it from this plot!

How do I plot these model results?


Rule for plotting model results:

- 1 If the model uses N variables, you should show all N effects *simultaneously*
- 2 If this is impractical, you should use a **partial effects plot**

Other names for partial effects:

Incorrect example, using raw data:

```
#Fit model with 5 variables (all important)
simMod <- lm(y~X1+X2+X3+X4+X5,data=pred)
#Incorrect way, using raw data
plot(y~X5,data=pred,pch=19,cex.lab=3)
```



The effect of $X5$ is actually **very** strong ($p < 0.0001$), but it doesn't look like it from this plot!

How do I plot these model results?

Rule for plotting model results:


- ① If the model uses N variables, you should show all N effects *simultaneously*
- ② If this is impractical, you should use a **partial effects plot**

Other names for partial effects:

- *counterfactual plot, predictor effect plot, leverage plot*

Incorrect example, using raw data:

```
#Fit model with 5 variables (all important)
simMod <- lm(y~X1+X2+X3+X4+X5,data=pred)
#Incorrect way, using raw data
plot(y~X5,data=pred,pch=19,cex.lab=3)
```



The effect of $X5$ is actually **very** strong ($p < 0.0001$), but it doesn't look like it from this plot!

How do I plot these model results?

Rule for plotting model results:


- ① If the model uses N variables, you should show all N effects *simultaneously*
- ② If this is impractical, you should use a **partial effects plot**

Other names for partial effects:

- *counterfactual* plot, *predictor effect* plot, *leverage* plot
- Try using `effects` or `ggeffects`. Requires the `effects` and `ggeffects` packages

Incorrect example, using raw data:

```
#Fit model with 5 variables (all important)
simMod <- lm(y~X1+X2+X3+X4+X5,data=pred)
#Incorrect way, using raw data
plot(y~X5,data=pred,pch=19,cex.lab=3)
```




The effect of $X5$ is actually **very** strong ($p < 0.0001$), but it doesn't look like it from this plot!


Partial effects plots - using *effects*

```
library(effects) #Load effects package
simModEff <- predictorEffects(simMod,partial.residuals=TRUE) #Calculate partial effects
#Plot partial effects
plot(simModEff,lines=list(col='red'), partial.residuals=list(pch=19,col='black',cex=0.25))
```


X1 predictor effect plot




X2 predictor effect plot




X3 predictor effect plot



X4 predictor effect plot




X5 predictor effect plot




Partial effects plots - using `sjmnredict`

```
library(ggeffects) #Load ggeffects package  
simModEff2 <- ggeffect(simMod,terms=c('X5')) #Calculate partial effects for X5  
plot(simModEff2) #Plot effect of X5
```




Interactions

What if the slopes *and* intercepts differ between groups?



Interactions



$$\begin{aligned} \hat{mpg} = & b_0 + b_1 disp \\ & + b_2 gears_4 + b_3 gears_5 \\ & + b_4 (disp \times gears_4) \\ & + b_5 (disp \times gears_5) \\ mpg \sim & Normal(\hat{mpg}, \sigma) \end{aligned}$$

- Interactions occur when predictors are *multiplied*
- In this case, *disp* is multiplied by *gears₄* and *gears₅*
- *gears* now changes the intercept and the slope of the relationship between *mpg* and *disp*

How do I get R to fit this model?

```
#mpg depends on disp interacted (*) with gears
mod2 <- lm(mpg ~ disp*factor(gear), data = mtcars)
summary(mod2)
```

```
##
## Call:
## lm(formula = mpg ~ disp * factor(gear), data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.5986 -1.5990 -0.0143  1.6329  4.9926
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    24.51556     2.462431   9.956 2.32e-10 ***
## disp          -0.025770     0.007265  -3.547 0.001505 **
## factor(gear)4    15.051963     3.558043   4.230 0.000256 ***
## factor(gear)5     7.145380     3.535913   2.021 0.053711 .
## disp:factor(gear)4 -0.096442     0.021261  -4.536 0.000114 ***
## disp:factor(gear)5 -0.025005     0.013320  -1.877 0.071742 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.579 on 26 degrees of freedom
## Multiple R-squared:  0.8465, Adjusted R-squared:  0.817
## F-statistic: 28.67 on 5 and 26 DF,  p-value: 8.452e-10
```

Beware of fitting too many interactions, or else the *Bilbo effect* occurs!

Dummy variables

```
mod2Matrix <- model.matrix(mod2) #Get model matrix (columns used to predict mpg)
colnames(mod2Matrix) <- gsub('factor\\(gear\\)', 'gear', colnames(mod2Matrix)) #Shorten colnames
head(mod2Matrix, 28) #Show first 28 rows of model matrix
```

```
##              (Intercept)  disp gear4 gear5 disp:gear4 disp:gear5
## Mazda RX4              1 160.0      1      0      160.0      0.0
## Mazda RX4 Wag          1 160.0      1      0      160.0      0.0
## Datsun 710              1 108.0      1      0      108.0      0.0
## Hornet 4 Drive          1 258.0      0      0       0.0      0.0
## Hornet Sportabout       1 360.0      0      0       0.0      0.0
## Valiant                 1 225.0      0      0       0.0      0.0
## Duster 360              1 360.0      0      0       0.0      0.0
## Merc 240D               1 146.7      1      0      146.7      0.0
## Merc 230                1 140.8      1      0      140.8      0.0
## Merc 280                1 167.6      1      0      167.6      0.0
## Merc 280C               1 167.6      1      0      167.6      0.0
## Merc 450SE              1 275.8      0      0       0.0      0.0
## Merc 450SL              1 275.8      0      0       0.0      0.0
## Merc 450SLC             1 275.8      0      0       0.0      0.0
## Cadillac Fleetwood      1 472.0      0      0       0.0      0.0
## Lincoln Continental      1 460.0      0      0       0.0      0.0
## Chrysler Imperial       1 440.0      0      0       0.0      0.0
## Fiat 128                 1  78.7      1      0       78.7      0.0
## Honda Civic              1  75.7      1      0       75.7      0.0
## Toyota Corolla           1  71.1      1      0       71.1      0.0
## Toyota Corona           1 120.1      0      0       0.0      0.0
## Dodge Challenger         1 318.0      0      0       0.0      0.0
## AMC Javelin              1 304.0      0      0       0.0      0.0
## Camaro Z28               1 350.0      0      0       0.0      0.0
## Pontiac Firebird         1 400.0      0      0       0.0      0.0
## Fiat X1-9                1  79.0      1      0       79.0      0.0
## Porsche 914-2           1 120.3      0      1       0.0     120.3
```

A challenger approaches!

- Since you're all bat folks, here's some bat data!

A challenger approaches!

- Since you're all bat folks, here's some bat data!
 - `batDat.csv`

A challenger approaches!

- Since you're all bat folks, here's some bat data!
 - `batDat.csv`
- Data: 100 bat weights from 2 cities, recorded along with sex and age

A challenger approaches!

- Since you're all bat folks, here's some bat data!
 - `batDat.csv`
- Data: 100 bat weights from 2 cities, recorded along with sex and age
- How do these variables affect bat weight?

A challenger approaches!

- Since you're all bat folks, here's some bat data!
 - `batDat.csv`
- Data: 100 bat weights from 2 cities, recorded along with sex and age
- How do these variables affect bat weight?
 - Think about how these variables might be related to weight using your brain

A challenger approaches!

- Since you're all bat folks, here's some bat data!
 - `batDat.csv`
- Data: 100 bat weights from 2 cities, recorded along with sex and age
- How do these variables affect bat weight?
 - Think about how these variables might be related to weight using your brain
 - Fit a model using `lm`

A challenger approaches!

- Since you're all bat folks, here's some bat data!
 - `batDat.csv`
- Data: 100 bat weights from 2 cities, recorded along with sex and age
- How do these variables affect bat weight?
 - Think about how these variables might be related to weight using your brain
 - Fit a model using `lm`
 - Make some plots, using `effects` or `ggeffects`

Part 3: Models behaving badly

Motivation

Are my model results reliable?

- Residual checks

Motivation

Are my model results reliable?

- Residual checks
- Transformations

Motivation

Are my model results reliable?


- Residual checks
- Transformations
- Collinearity

Motivation

Are my model results reliable?

- Residual checks
- Transformations
- Collinearity
- How much stuff should I put into my model?

Assumptions of linear regression




$$\hat{mpg} = b_0 + b_1 disp$$

There are 3 main assumptions to this model:

- ① The relationship between *disp* and *mpg* is linear

This is pretty easy to see if you only have 1 variable, but...

Assumptions of linear regression




$$\hat{mpg} = b_0 + b_1 disp$$

There are 3 main assumptions to this model:

- ① The relationship between *disp* and *mpg* is linear
- ② *mpg* (the data) is Normally distributed around \hat{mpg} (the line)

This is pretty easy to see if you only have 1 variable, but...

Assumptions of linear regression




$$\hat{mpg} = b_0 + b_1 disp$$

There are 3 main assumptions to this model:

- ① The relationship between *disp* and *mpg* is linear
- ② *mpg* (the data) is Normally distributed around \hat{mpg} (the line)
- ③ σ is the same everywhere

This is pretty easy to see if you only have 1 variable, but...

What if I have many variables?





Difficult to see if the assumptions are met

Solution: residual checks

Some common ways of checking the assumptions: **residual plots**



```
mod1 <- lm(mpg~disp*factor(gear),data=mtcars) #Fits model
par(mfrow=c(1,2),mar=c(3,3,1,1)+1) #Splits plot into 2
plot(mod1, which=c(1,2)) #1st and 2nd residual plots
```




Solution: residual checks

Some common ways of checking the assumptions: **residual plots**

```
mod1 <- lm(mpg~disp*factor(gear),data=mtcars) #Fits model
par(mfrow=c(1,2),mar=c(3,3,1,1)+1) #Splits plot into 2
plot(mod1, which=c(1,2)) #1st and 2nd residual plots
```




Problem 1: Non-linear relationship




```
lm(y1~x,data=d1)
```

y_1 clearly follows a hump-shaped relationship, not a linear one.




Solution: transform predictors




```
lm(y1~poly(x,2),data=d1)
```

log and *square-root* transformations are common




Problem 2a: Non-normal response




```
lm(y2~x,data=d1)
```

y_2 is count data (integers ≥ 0). Very common in ecological data.




Solution: transform data to meet assumptions




```
lm(log(y2+1)~x,data=d1)
```

Square-root transformations are also common




Problem 2b: Non-normal response




```
lm(y3~x,data=d1)
```

y_3 is binomial data (success/failure, 0 or 1). Very common in ecological data.




Solution: use a Generalized Linear Model (GLM)





• This is a topic for another lecture. Hold tight!

Problem: variables are on different scales




```
lm(y4~x3,data=d1)
```

- y_4 is tiny, while x_3 is huge





Problem: variables are on different scales




```
lm(y4~x3,data=d1)
```


- y_4 is tiny, while x_3 is huge



Solution: scale data/predictors before fitting




```
#Subtracts mean, divides by SD  
d1$s.y4 <- scale(y4)  
d1$s.x3 <- scale(x3)  
lm(s.y4~s.x3,data=d1) #Refit
```




- Residuals are the same as before

Solution: scale data/predictors before fitting



```
#Subtracts mean, divides by SD  
d1$s.y4 <- scale(y4)  
d1$s.x3 <- scale(x3)  
lm(s.y4~s.x3,data=d1) #Refit
```





- Residuals are the same as before
- Coefficients are now related to *scaled* data and predictor

But wait... there's more (assumptions)!

One more assumption:

- ④ If you have 2+ predictors in your model, the predictors are not related to each other

```
lm(y0~x+x2,data=d1)
```





But wait... there's more (assumptions)!

One more assumption:

- ④ If you have 2+ predictors in your model, the predictors are not related to each other
 - Say we have 2 predictors, x and x_2 :

```
lm(y0~x+x2,data=d1)
```



But wait... there's more (assumptions)!



One more assumption:

- ④ If you have 2+ predictors in your model, the predictors are not related to each other

- Say we have 2 predictors, x and x_2 :

```
lm(y0~x+x2,data=d1)
```

- Model fits, and residuals look OK, but there's trouble ahead!




Uh oh! Collinearity!

```
#Function to print correlation (r) value
```

```
corText <- function(x,y){  
  text(0.5,0.5,round(cor(x,y),3))  
}
```

```
#Pairplot of y0, x, and x2
```

```
pairs(d1[,c('y0','x','x2')],lower.panel=corText)
```



- x and x2 mean basically the same thing!

```
library(car)
```

```
#VIF scores:
```

```
# 1 = no problem
```

```
# 1-5 = some problems
```

```
# 5+ = big problems!
```

```
vif(m2)
```

```
##          x          x2  
## 11.31812 11.31812
```

pairs() is useful for looking at relations


Uh oh! Collinearity!

```
#Function to print correlation (r) value
```

```
corText <- function(x,y){  
  text(0.5,0.5,round(cor(x,y),3))  
}
```

```
#Pairplot of y0, x, and x2
```

```
pairs(d1[,c('y0','x','x2')],lower.panel=corText)
```



- x and x2 mean basically the same thing!
- Also revealed using variance-inflation factors (VIFs):

```
library(car)
```

```
#VIF scores:
```

```
# 1 = no problem
```

```
# 1-5 = some problems
```

```
# 5+ = big problems!
```

```
vif(m2)
```

```
##          x          x2  
## 11.31812 11.31812
```

pairs() is useful for looking at relations

Is collinearity really that bad?

#Correct model

```
m1 <- lm(y0~x,data=d1)
```

	Estimate	Std. Error	Pr(> t)
(Intercept)	0.7851936	0.1943002	0.0001059
x	-0.1900346	0.0342596	0.0000002

#Incorrect model

```
m2 <- lm(y0~x+x2,data=d1)
```

	Estimate	Std. Error	Pr(> t)
(Intercept)	0.7860300	0.1955770	0.0001155
x	-0.1812556	0.1158464	0.1209288
x2	-0.0094931	0.1196074	0.9369028

- Increases SE of each term, so model may “miss” important terms

Is collinearity really that bad?

#Correct model

```
m1 <- lm(y0~x,data=d1)
```

	Estimate	Std. Error	Pr(> t)
(Intercept)	0.7851936	0.1943002	0.0001059
x	-0.1900346	0.0342596	0.0000002

#Incorrect model

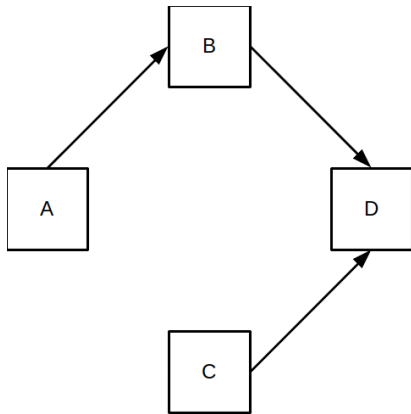
```
m2 <- lm(y0~x+x2,data=d1)
```

	Estimate	Std. Error	Pr(> t)
(Intercept)	0.7860300	0.1955770	0.0001155
x	-0.1812556	0.1158464	0.1209288
x2	-0.0094931	0.1196074	0.9369028

- Increases SE of each term, so model may “miss” important terms
- Gets worse with increasing correlation, or if many terms are correlated!

How do we fix this? Depends on your goals:


- 1 I care about predicting things



`lm(D ~ B + C)`

How do we fix this? Depends on your goals:


- 1 I care about predicting things
 - Use dimensional reduction (e.g. PCA) and re-run model



$\text{lm}(D \sim B + C)$

How do we fix this? Depends on your goals:


- 1 I care about predicting things
 - Use dimensional reduction (e.g. PCA) and re-run model
- 2 I care about what's causing things



$$\text{lm}(D \sim B + C)$$

How do we fix this? Depends on your goals:


- ① I care about predicting things
 - Use dimensional reduction (e.g. PCA) and re-run model
- ② I care about what's causing things
 - Design experiment to separate cause and effect



$$\text{lm}(D \sim B + C)$$

How do we fix this? Depends on your goals:

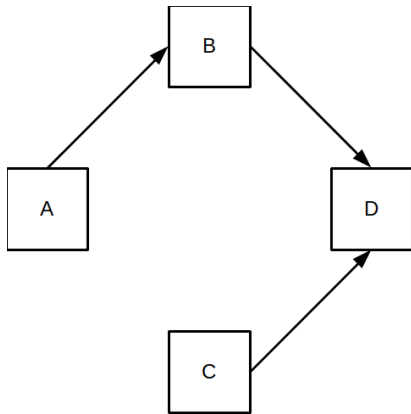
- ① I care about predicting things
 - Use dimensional reduction (e.g. PCA) and re-run model
- ② I care about what's causing things
 - Design experiment to separate cause and effect
 - Think about what is causing what.
Graphical models are helpful for this



$$\text{lm}(D \sim B + C)$$

How do we fix this? Depends on your goals:


- ① I care about predicting things
 - Use dimensional reduction (e.g. PCA) and re-run model
- ② I care about what's causing things
 - Design experiment to separate cause and effect
 - Think about what is causing what.
Graphical models are helpful for this
 - Not all variables have to be included!



$$\text{lm}(D \sim B + C)$$

How do we fix this? Depends on your goals:

- ① I care about predicting things
 - Use dimensional reduction (e.g. PCA) and re-run model
- ② I care about what's causing things
 - Design experiment to separate cause and effect
 - Think about what is causing what.
Graphical models are helpful for this
 - Not all variables have to be included!



- Simple graphical model, where the effect of A on D is *mediated* by B.


$\text{lm}(D \sim B + C)$

How do we fix this? Depends on your goals:

- ① I care about predicting things
 - Use dimensional reduction (e.g. PCA) and re-run model
- ② I care about what's causing things
 - Design experiment to separate cause and effect
 - Think about what is causing what.

Graphical models are helpful for this

- Not all variables have to be included!



- Simple graphical model, where the effect of A on D is *mediated* by B.
- “Correct” lm model of D:

$$\text{lm}(D \sim B + C)$$

A challenger approaches!

- Guess what... more bat data! This time there are 6 variables that were measured. We're interested in predicting bats (counts of bats per night).

A challenger approaches!

- Guess what... more bat data! This time there are 6 variables that were measured. We're interested in predicting bats (counts of bats per night).
- Formulate a causal model that seems reasonable

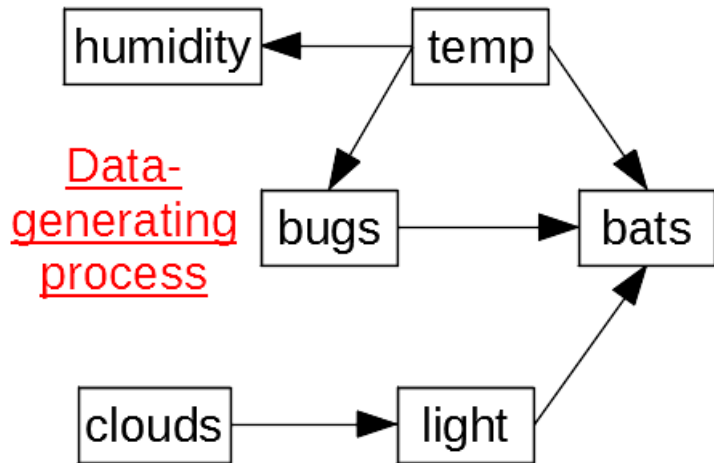
A challenger approaches!

- Guess what... more bat data! This time there are 6 variables that were measured. We're interested in predicting bats (counts of bats per night).
- Formulate a causal model that seems reasonable
 - Draw it out on paper/in PowerPoint using flow diagrams

A challenger approaches!

- Guess what... more bat data! This time there are 6 variables that were measured. We're interested in predicting bats (counts of bats per night).
- Formulate a causal model that seems reasonable
 - Draw it out on paper/in PowerPoint using flow diagrams
- Fit an `lm` model of bats from your causal model, check the assumptions, and update as necessary

Here's the answer



This is the **true** process that generated the data. Model for bats should look like:

```
lm(log(bats+0.1)~poly(temp,2)+light+bugs,data=dat)
```