

Mixed effects models

“Wheels within wheels”


Samuel Robinson, Ph.D.

Oct 6, 2023

Part 1: Mixed effects models


Outline

- Linear mixed effects models (LMMs)




Outline

- Linear mixed effects models (LMMs)
 - A bit of math




Outline

- Linear mixed effects models (LMMs)
 - A bit of math
 - Fixed vs. random effects




Outline

- Linear mixed effects models (LMMs)
 - A bit of math
 - Fixed vs. random effects
 - Random intercepts and slopes




Outline

- Linear mixed effects models (LMMs)
 - A bit of math
 - Fixed vs. random effects
 - Random intercepts and slopes
- Generalized linear mixed effects models (GLMMs)




Outline

- Linear mixed effects models (LMMs)
 - A bit of math
 - Fixed vs. random effects
 - Random intercepts and slopes
- Generalized linear mixed effects models (GLMMs)
 - Residuals checks




Outline

- Linear mixed effects models (LMMs)
 - A bit of math
 - Fixed vs. random effects
 - Random intercepts and slopes
- Generalized linear mixed effects models (GLMMs)
 - Residuals checks
 - Some sage advice




Outline

- Linear mixed effects models (LMMs)
 - A bit of math
 - Fixed vs. random effects
 - Random intercepts and slopes
- Generalized linear mixed effects models (GLMMs)
 - Residuals checks
 - Some sage advice
- Hypothesis testing and inference




Outline

- Linear mixed effects models (LMMs)
 - A bit of math
 - Fixed vs. random effects
 - Random intercepts and slopes
- Generalized linear mixed effects models (GLMMs)
 - Residuals checks
 - Some sage advice
- Hypothesis testing and inference
 - Slopes and intercepts




Outline

- Linear mixed effects models (LMMs)
 - A bit of math
 - Fixed vs. random effects
 - Random intercepts and slopes
- Generalized linear mixed effects models (GLMMs)
 - Residuals checks
 - Some sage advice
- Hypothesis testing and inference
 - Slopes and intercepts
 - Entire terms



Outline

- Linear mixed effects models (LMMs)
 - A bit of math
 - Fixed vs. random effects
 - Random intercepts and slopes
- Generalized linear mixed effects models (GLMMs)
 - Residuals checks
 - Some sage advice
- Hypothesis testing and inference
 - Slopes and intercepts
 - Entire terms
 - AIC, R^2 , and Star-Gazing



What are mixed effects models?

Many different names:

- 1 Mixed effects models

I usually use the term *heirarchical models*, as this is the closest to what I will teach you

What are mixed effects models?

Many different names:

- ① Mixed effects models
- ② Random effects models

I usually use the term *heirarchical models*, as this is the closest to what I will teach you

What are mixed effects models?

Many different names:

- ① Mixed effects models
- ② Random effects models
- ③ Hierarchical models

I usually use the term *heirarchical models*, as this is the closest to what I will teach you

What are mixed effects models?

Many different names:

- ① Mixed effects models
- ② Random effects models
- ③ Hierarchical models
- ④ Empirical/Bayesian hierarchical models

I usually use the term *heirarchical models*, as this is the closest to what I will teach you

What are mixed effects models?

Many different names:

- ① Mixed effects models
- ② Random effects models
- ③ Hierarchical models
- ④ Empirical/Bayesian hierarchical models
- ⑤ Latent variable models

I usually use the term *heirarchical models*, as this is the closest to what I will teach you

What are mixed effects models?

Many different names:

- ① Mixed effects models
- ② Random effects models
- ③ Hierarchical models
- ④ Empirical/Bayesian hierarchical models
- ⑤ Latent variable models
- ⑥ Split-plot models

I usually use the term *heirarchical models*, as this is the closest to what I will teach you

What are mixed effects models?

Many different names:

- ① Mixed effects models
- ② Random effects models
- ③ Hierarchical models
- ④ Empirical/Bayesian hierarchical models
- ⑤ Latent variable models
- ⑥ Split-plot models
- ⑦ Variance partitioning

I usually use the term *heirarchical models*, as this is the closest to what I will teach you

Scary math

Unfortunately, we need a review of matrix algebra in order to explain this:

- This is a matrix:

$$A = \begin{bmatrix} 1 & 4 & 7 \\ 2 & 5 & 8 \\ 3 & 6 & 9 \end{bmatrix}$$

$$A \times b = Ab = 1 \times \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} + 2 \times \begin{bmatrix} 4 \\ 5 \\ 6 \end{bmatrix} + 3 \times \begin{bmatrix} 7 \\ 8 \\ 9 \end{bmatrix} = \begin{bmatrix} 30 \\ 36 \\ 42 \end{bmatrix}$$

Scary math

Unfortunately, we need a review of matrix algebra in order to explain this:

- This is a matrix:

$$A = \begin{bmatrix} 1 & 4 & 7 \\ 2 & 5 & 8 \\ 3 & 6 & 9 \end{bmatrix}$$

- This is a vector:

$$b = \begin{bmatrix} 1 & 2 & 3 \end{bmatrix}$$

$$A \times b = Ab = 1 \times \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} + 2 \times \begin{bmatrix} 4 \\ 5 \\ 6 \end{bmatrix} + 3 \times \begin{bmatrix} 7 \\ 8 \\ 9 \end{bmatrix} = \begin{bmatrix} 30 \\ 36 \\ 42 \end{bmatrix}$$

Scary math

Unfortunately, we need a review of matrix algebra in order to explain this:

- This is a matrix:

$$A = \begin{bmatrix} 1 & 4 & 7 \\ 2 & 5 & 8 \\ 3 & 6 & 9 \end{bmatrix}$$

- This is a vector:

$$b = \begin{bmatrix} 1 & 2 & 3 \end{bmatrix}$$

- Multiplying them looks like this:

$$A \times b = Ab = 1 \times \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} + 2 \times \begin{bmatrix} 4 \\ 5 \\ 6 \end{bmatrix} + 3 \times \begin{bmatrix} 7 \\ 8 \\ 9 \end{bmatrix} = \begin{bmatrix} 30 \\ 36 \\ 42 \end{bmatrix}$$

Why do we call them “linear models”?

- *Linear* mapping of **coefficients** onto a **model matrix** (from your data)

Why do we call them “linear models”?

- *Linear* mapping of **coefficients** onto a **model matrix** (from your data)
- Coefficients:

$$\beta = \begin{bmatrix} 0.1 & 1.8 & -0.03 \end{bmatrix}$$

Why do we call them “linear models”?

- *Linear* mapping of **coefficients** onto a **model matrix** (from your data)
- Coefficients:

$$\beta = \begin{bmatrix} 0.1 & 1.8 & -0.03 \end{bmatrix}$$

- Model matrix:

$$X = \begin{bmatrix} 1 & 1 & 10 \\ 1 & 1 & 12 \\ 1 & 0 & 9 \\ \vdots & \vdots & \vdots \end{bmatrix}$$

Why do we call them “linear models”?

- *Linear* mapping of **coefficients** onto a **model matrix** (from your data)

- Coefficients:

$$\beta = \begin{bmatrix} 0.1 & 1.8 & -0.03 \end{bmatrix}$$

- Model matrix:

$$X = \begin{bmatrix} 1 & 1 & 10 \\ 1 & 1 & 12 \\ 1 & 0 & 9 \\ \vdots & \vdots & \vdots \end{bmatrix}$$

- Multiplying them looks like:

$$\hat{y} = X\beta = \begin{bmatrix} 1.60 \\ 1.54 \\ -0.17 \\ \vdots \end{bmatrix}$$

This is exactly what R does to fit models:

```
head(dat)
```

```
##           y           x site
## 1  1.5101095 -4.248450    g
## 2  3.7190900  5.766103    j
## 3 -4.3737644 -1.820462    f
## 4 30.1459331  7.660348    n
## 5  0.2777422  8.809346    o
## 6 -3.6978175 -9.088870    p
```

```
m1 <- lm(y~x,data=dat) #Uses x to predict y
summary(m1)
```

```
##
## Call:
## lm(formula = y ~ x, data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -18.2574  -4.1262   0.0296   3.1854  25.2780
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.7306     0.5772   4.731 4.92e-06 ***
## x              0.7213     0.1020   7.074 4.60e-11 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.3 on 158 degrees of freedom
## Multiple R-squared:  0.2405, Adjusted R-squared:  0.2357
## F-statistic: 50.04 on 1 and 158 DF,  p-value: 4.6e-11
```

This is exactly what R does to fit models (cont.):

```
head(model.matrix(m1))
```

```
##      (Intercept)          x
## 1             1 -4.248450
## 2             1  5.766103
## 3             1 -1.820462
## 4             1  7.660348
## 5             1  8.809346
## 6             1 -9.088870
```

```
coef(m1)
```

```
##      (Intercept)          x
## 2.7305689  0.7212867
```

```
pred2 <- model.matrix(m1) %*% coef(m1) #predicted = matrix * coefs
head(data.frame(pred1=predict(m1),pred2)) #same thing!
```

```
##      pred1      pred2
## 1 -0.3337812 -0.3337812
## 2  6.8895819  6.8895819
## 3  1.4174942  1.4174942
## 4  8.2558758  8.2558758
## 5  9.0846325  9.0846325
## 6 -3.8251119 -3.8251119
```

Groups are coded by “dummy variables” (0s and 1s)

```
m2 <- lm(y~site,data=dat) #Use site to predict y  
head(model.matrix(m2)) #0s and 1s used to identify groups
```

```
## (Intercept) siteb sitec sited sitee sitef siteg siteh sitei sitej sitek sitel  
## 1          1      0      0      0      0      0      1      0      0      0      0  
## 2          1      0      0      0      0      0      0      0      0      1      0  
## 3          1      0      0      0      0      1      0      0      0      0      0  
## 4          1      0      0      0      0      0      0      0      0      0      0  
## 5          1      0      0      0      0      0      0      0      0      0      0  
## 6          1      0      0      0      0      0      0      0      0      0      0  
##      sitem siten siteo sitep  
## 1      0      0      0      0  
## 2      0      0      0      0  
## 3      0      0      0      0  
## 4      0      1      0      0  
## 5      0      0      1      0  
## 6      0      0      0      1
```

```
coef(m2) #This uses the 1st site as the "control" group
```

```
## (Intercept)      siteb      sitec      sited      sitee      sitef  
##  7.192416 -11.998464 -14.632803   1.983649  -7.765354  -4.523079  
##      siteg      siteh      sitei      sitej      sitek      sitel  
## -3.439621  -8.280601  -4.306456  -4.085855  -5.663021  -5.155112  
##      sitem      siten      siteo      sitep  
## -6.226642   8.403599  -8.626661 -10.934182
```

Structure of LMs... now with matrices!

- All linear models take the form:

$$\hat{y} = X\beta = b_0 1 + b_1 x_1 \dots + b_i x_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

Structure of LMs... now with matrices!

- All linear models take the form:

$$\hat{y} = X\beta = b_0\mathbf{1} + b_1x_1 \dots + b_ix_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is a vector of data you want to predict

Structure of LMs... now with matrices!

- All linear models take the form:

$$\hat{y} = X\beta = b_0\mathbf{1} + b_1x_1 \dots + b_ix_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is a vector of data you want to predict
- \hat{y} is a vector of *predicted values* for y

Structure of LMs... now with matrices!

- All linear models take the form:

$$\hat{y} = X\beta = b_0\mathbf{1} + b_1x_1\ldots + b_ix_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is a vector of data you want to predict
- \hat{y} is a vector of *predicted values* for y
- $X = \{1, x_1\ldots\}$ is a matrix of *predictors* for y

Structure of LMs... now with matrices!

- All linear models take the form:

$$\hat{y} = X\beta = b_0\mathbf{1} + b_1x_1\ldots + b_ix_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is a vector of data you want to predict
- \hat{y} is a vector of *predicted values* for y
- $X = \{1, x_1\ldots\}$ is a matrix of *predictors* for y
- $\beta = \{b_0, b_1, \ldots\}$ is a vector of *coefficients*

Structure of LMs... now with matrices!

- All linear models take the form:

$$\hat{y} = X\beta = b_0 1 + b_1 x_1 \dots + b_i x_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is a vector of data you want to predict
- \hat{y} is a vector of *predicted values* for y
- $X = \{1, x_1 \dots\}$ is a matrix of *predictors* for y
- $\beta = \{b_0, b_1, \dots\}$ is a vector of *coefficients*
- $y \sim \text{Normal}(\hat{y}, \sigma)$ means:

Structure of LMs... now with matrices!

- All linear models take the form:

$$\hat{y} = X\beta = b_0 1 + b_1 x_1 \dots + b_i x_i$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- y is a vector of data you want to predict
- \hat{y} is a vector of *predicted values* for y
- $X = \{1, x_1 \dots\}$ is a matrix of *predictors* for y
- $\beta = \{b_0, b_1, \dots\}$ is a vector of *coefficients*
- $y \sim \text{Normal}(\hat{y}, \sigma)$ means:
 - “ y follows a Normal distribution with mean \hat{y} and SD σ ”

Fixed effects vs. Random effects

Say that X is a model matrix coding for 10 sites¹, and y is something we're interested in predicting

$$\hat{y} = b_0 + X\beta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\hat{y} = b_0 + X\zeta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$
$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- Site coefficients (β) are unrelated to each other

¹Intercept is a separate variable

Fixed effects vs. Random effects

Say that X is a model matrix coding for 10 sites¹, and y is something we're interested in predicting

$$\hat{y} = b_0 + X\beta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\hat{y} = b_0 + X\zeta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$
$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- Site coefficients (β) are unrelated to each other
- σ is the SD of *residuals*

¹Intercept is a separate variable

Fixed effects vs. Random effects

Say that X is a model matrix coding for 10 sites¹, and y is something we're interested in predicting

$$\hat{y} = b_0 + X\beta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\hat{y} = b_0 + X\zeta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$
$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- Site coefficients (β) are unrelated to each other
- σ is the SD of *residuals*
- Site is a **fixed effect**

¹Intercept is a separate variable

Fixed effects vs. Random effects

Say that X is a model matrix coding for 10 sites¹, and y is something we're interested in predicting

$$\hat{y} = b_0 + X\beta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- Site coefficients (β) are unrelated to each other
- σ is the SD of *residuals*
- Site is a **fixed effect**

$$\hat{y} = b_0 + X\zeta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$
$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- Site coefficients (ζ) are related to each other via a *Normal* distribution

¹Intercept is a separate variable

Fixed effects vs. Random effects

Say that X is a model matrix coding for 10 sites¹, and y is something we're interested in predicting

$$\hat{y} = b_0 + X\beta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- Site coefficients (β) are unrelated to each other
- σ is the SD of *residuals*
- Site is a **fixed effect**

$$\hat{y} = b_0 + X\zeta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$
$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- Site coefficients (ζ) are related to each other via a *Normal* distribution
- σ is the SD of *residuals*, σ_{site} is the SD of *sites*

¹Intercept is a separate variable

Fixed effects vs. Random effects

Say that X is a model matrix coding for 10 sites¹, and y is something we're interested in predicting

$$\hat{y} = b_0 + X\beta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- Site coefficients (β) are unrelated to each other
- σ is the SD of *residuals*
- Site is a **fixed effect**

$$\hat{y} = b_0 + X\zeta$$
$$y \sim \text{Normal}(\hat{y}, \sigma)$$
$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- Site coefficients (ζ) are related to each other via a *Normal* distribution
- σ is the SD of *residuals*, σ_{site} is the SD of *sites*
- Site is a **random effect**

¹Intercept is a separate variable

Mixed effects = fixed + random effects

A mixed effects model has both **fixed** and **random** effects

$$\hat{y} = X\beta + U\zeta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- X = fixed effects matrix (e.g. intercept, temperature)

Mixed effects = fixed + random effects

A mixed effects model has both **fixed** and **random** effects

$$\hat{y} = X\beta + U\zeta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- X = fixed effects matrix (e.g. intercept, temperature)
- β = fixed effects coefficients

Mixed effects = fixed + random effects

A mixed effects model has both **fixed** and **random** effects

$$\hat{y} = X\beta + U\zeta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- X = fixed effects matrix (e.g. intercept, temperature)
- β = fixed effects coefficients
- U = random effects matrix (e.g. sites)

Mixed effects = fixed + random effects

A mixed effects model has both **fixed** and **random** effects

$$\hat{y} = X\beta + U\zeta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- X = fixed effects matrix (e.g. intercept, temperature)
- β = fixed effects coefficients
- U = random effects matrix (e.g. sites)
- ζ = random effects coefficients

Mixed effects = fixed + random effects

A mixed effects model has both **fixed** and **random** effects

$$\hat{y} = X\beta + U\zeta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- X = fixed effects matrix (e.g. intercept, temperature)
- β = fixed effects coefficients
- U = random effects matrix (e.g. sites)
- ζ = random effects coefficients
- $\sigma, \sigma_{\text{site}}$ = variance terms

Mixed effect model example

Let's go back to our earlier example:

- We're interested in predicting y using x (fixed effects)

```
head(dat)
```

```
##           y           x site
## 1  1.5101095 -4.248450    g
## 2  3.7190900  5.766103    j
## 3 -4.3737644 -1.820462    f
## 4 30.1459331  7.660348    n
## 5  0.2777422  8.809346    o
## 6 -3.6978175 -9.088870    p
```

```
library(lme4) #Mixed effects library
mm1 <- lmer(y ~ x + (1|site),data=dat) #site is fit as "random intercepts"
```

Mixed effect model example

Let's go back to our earlier example:

- We're interested in predicting y using x (fixed effects)
- Data was collected at a number of *sites*, which may affect y "somehow"

```
head(dat)
```

```
##           y           x site
## 1  1.5101095 -4.248450    g
## 2  3.7190900  5.766103    j
## 3 -4.3737644 -1.820462    f
## 4 30.1459331  7.660348    n
## 5  0.2777422  8.809346    o
## 6 -3.6978175 -9.088870    p
```

```
library(lme4) #Mixed effects library
mm1 <- lmer(y ~ x + (1|site),data=dat) #site is fit as "random intercepts"
```

Mixed effect model example

Let's go back to our earlier example:

- We're interested in predicting y using x (fixed effects)
- Data was collected at a number of *sites*, which may affect y "somehow"
- Effect of each site is normally distributed

```
head(dat)
```

```
##           y           x site
## 1  1.5101095 -4.248450    g
## 2  3.7190900  5.766103    j
## 3 -4.3737644 -1.820462    f
## 4 30.1459331  7.660348    n
## 5  0.2777422  8.809346    o
## 6 -3.6978175 -9.088870    p
```

```
library(lme4) #Mixed effects library
mm1 <- lmer(y ~ x + (1|site),data=dat) #site is fit as "random intercepts"
```

Mixed effect model example

```
summary(mm1)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ x + (1 | site)
## Data: dat
##
## REML criterion at convergence: 1040.4
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.50816 -0.71380 -0.02682  0.69401  3.01951
##
## Random effects:
## Groups Name Variance Std.Dev.
## site (Intercept) 26.78 5.175
## Residual 31.68 5.628
## Number of obs: 160, groups: site, 16
##
## Fixed effects:
## Estimate Std. Error t value
## (Intercept) 2.00331 1.38297 1.449
## x 0.76192 0.08083 9.426
##
## Correlation of Fixed Effects:
## (Intr)
## x -0.002
```

Results from lmer model:

- Random effects:

²'lme4' author doesn't think they can be calculated. I somewhat agree

Mixed effect model example

```
summary(mm1)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ x + (1 | site)
## Data: dat
##
## REML criterion at convergence: 1040.4
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.50816 -0.71380 -0.02682  0.69401  3.01951
##
## Random effects:
## Groups Name Variance Std.Dev.
## site (Intercept) 26.78 5.175
## Residual 31.68 5.628
## Number of obs: 160, groups: site, 16
##
## Fixed effects:
## Estimate Std. Error t value
## (Intercept) 2.00331 1.38297 1.449
## x 0.76192 0.08083 9.426
##
## Correlation of Fixed Effects:
## (Intr)
## x -0.002
```

Results from lmer model:

- Random effects:
 - *residual* and *site* variance (σ , σ_{site})

²'lme4' author doesn't think they can be calculated. I somewhat agree

Mixed effect model example

```
summary(mm1)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ x + (1 | site)
## Data: dat
##
## REML criterion at convergence: 1040.4
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.50816 -0.71380 -0.02682  0.69401  3.01951
##
## Random effects:
## Groups Name Variance Std.Dev.
## site (Intercept) 26.78 5.175
## Residual 31.68 5.628
## Number of obs: 160, groups: site, 16
##
## Fixed effects:
## Estimate Std. Error t value
## (Intercept) 2.00331 1.38297 1.449
## x 0.76192 0.08083 9.426
##
## Correlation of Fixed Effects:
## (Intr)
## x -0.002
```

Results from lmer model:

- Random effects:
 - *residual* and *site* variance (σ , σ_{site})
- Fixed effects:

²'lme4' author doesn't think they can be calculated. I somewhat agree

Mixed effect model example

```
summary(mm1)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ x + (1 | site)
## Data: dat
##
## REML criterion at convergence: 1040.4
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.50816 -0.71380 -0.02682  0.69401  3.01951
##
## Random effects:
## Groups Name Variance Std.Dev.
## site (Intercept) 26.78 5.175
## Residual 31.68 5.628
## Number of obs: 160, groups: site, 16
##
## Fixed effects:
## Estimate Std. Error t value
## (Intercept) 2.00331 1.38297 1.449
## x 0.76192 0.08083 9.426
##
## Correlation of Fixed Effects:
## (Intr)
## x -0.002
```

Results from lmer model:

- Random effects:
 - *residual* and *site* variance (σ , σ_{site})
- Fixed effects:
 - Intercept and slope estimates (β)

²'lme4' author doesn't think they can be calculated. I somewhat agree

Mixed effect model example

```
summary(mm1)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ x + (1 | site)
## Data: dat
##
## REML criterion at convergence: 1040.4
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.50816 -0.71380 -0.02682  0.69401  3.01951
##
## Random effects:
## Groups Name Variance Std.Dev.
## site (Intercept) 26.78 5.175
## Residual 31.68 5.628
## Number of obs: 160, groups: site, 16
##
## Fixed effects:
## Estimate Std. Error t value
## (Intercept) 2.00331 1.38297 1.449
## x 0.76192 0.08083 9.426
##
## Correlation of Fixed Effects:
## (Intr)
## x -0.002
```

Results from lmer model:

- Random effects:
 - *residual* and *site* variance (σ , σ_{site})
- Fixed effects:
 - Intercept and slope estimates (β)
 - No d.f. and p-value²

²'lme4' author doesn't think they can be calculated. I somewhat agree

Mixed effect model example

```
summary(mm1)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ x + (1 | site)
## Data: dat
##
## REML criterion at convergence: 1040.4
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.50816 -0.71380 -0.02682  0.69401  3.01951
##
## Random effects:
## Groups Name Variance Std.Dev.
## site (Intercept) 26.78 5.175
## Residual 31.68 5.628
## Number of obs: 160, groups: site, 16
##
## Fixed effects:
## Estimate Std. Error t value
## (Intercept) 2.00331 1.38297 1.449
## x 0.76192 0.08083 9.426
##
## Correlation of Fixed Effects:
## (Intr)
## x -0.002
```


Results from lmer model:

- Random effects:
 - *residual* and *site* variance (σ , σ_{site})
- Fixed effects:
 - Intercept and slope estimates (β)
 - No d.f. and p-value²
 - If you need p-values for parameters, you can use the *lmerTest* package (or just calculate them yourself using means/SEs)

²'lme4' author doesn't think they can be calculated. I somewhat agree


Mixed effect model results

- In a *random intercepts* model, the regression line of x on y is allowed to move up or down around the main regression line for each site



Mixed effect model results



- In a *random intercepts* model, the regression line of x on y is allowed to move up or down around the main regression line for each site
- These changes in intercepts are *normally distributed*



Mixed effect model results (cont.)

For plotting, we want a partial effects plot that marginalizes across sites (i.e. “What does the trend look like at the average site?”)

- `ggpredict` works well for this. If you want partial residuals, you'll have to add them in yourself using `predict` and `residual`



Random slopes + intercepts

Suppose that y wasn't just higher or lower at each site, but that the effect of x on y was higher or lower at each site

$$\hat{y} = X\beta + U\zeta_{int} + U_x\zeta_{slope}$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta_{int} \sim \text{Normal}(0, \sigma_{int})$$

$$\zeta_{slope} \sim \text{Normal}(0, \sigma_{slope})$$

- X = fixed effects matrix (e.g. intercept, temperature)

Random slopes + intercepts

Suppose that y wasn't just higher or lower at each site, but that the effect of x on y was higher or lower at each site

$$\hat{y} = X\beta + U\zeta_{int} + U_x\zeta_{slope}$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta_{int} \sim \text{Normal}(0, \sigma_{int})$$

$$\zeta_{slope} \sim \text{Normal}(0, \sigma_{slope})$$

- X = fixed effects matrix (e.g. intercept, temperature)
- β = fixed effects coefficients

Random slopes + intercepts

Suppose that y wasn't just higher or lower at each site, but that the effect of x on y was higher or lower at each site

$$\hat{y} = X\beta + U\zeta_{int} + U_x\zeta_{slope}$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta_{int} \sim \text{Normal}(0, \sigma_{int})$$

$$\zeta_{slope} \sim \text{Normal}(0, \sigma_{slope})$$

- X = fixed effects matrix (e.g. intercept, temperature)
- β = fixed effects coefficients
- U = random intercept matrix (e.g. sites)

Random slopes + intercepts

Suppose that y wasn't just higher or lower at each site, but that the effect of x on y was higher or lower at each site

$$\hat{y} = X\beta + U\zeta_{int} + U_x\zeta_{slope}$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta_{int} \sim \text{Normal}(0, \sigma_{int})$$

$$\zeta_{slope} \sim \text{Normal}(0, \sigma_{slope})$$

- X = fixed effects matrix (e.g. intercept, temperature)
- β = fixed effects coefficients
- U = random intercept matrix (e.g. sites)
- U_x = random slopes matrix (e.g. temperature)

Random slopes + intercepts

Suppose that y wasn't just higher or lower at each site, but that the effect of x on y was higher or lower at each site

$$\hat{y} = X\beta + U\zeta_{int} + U_x\zeta_{slope}$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta_{int} \sim \text{Normal}(0, \sigma_{int})$$

$$\zeta_{slope} \sim \text{Normal}(0, \sigma_{slope})$$

- X = fixed effects matrix (e.g. intercept, temperature)
- β = fixed effects coefficients
- U = random intercept matrix (e.g. sites)
- U_x = random slopes matrix (e.g. temperature)
- $\zeta_{int}, \zeta_{slope}$ = random intercept and slope coefficients

Random slopes + intercepts

Suppose that y wasn't just higher or lower at each site, but that the effect of x on y was higher or lower at each site

$$\hat{y} = X\beta + U\zeta_{int} + U_x\zeta_{slope}$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta_{int} \sim \text{Normal}(0, \sigma_{int})$$

$$\zeta_{slope} \sim \text{Normal}(0, \sigma_{slope})$$

- X = fixed effects matrix (e.g. intercept, temperature)
- β = fixed effects coefficients
- U = random intercept matrix (e.g. sites)
- U_x = random slopes matrix (e.g. temperature)
- $\zeta_{int}, \zeta_{slope}$ = random intercept and slope coefficients
- $\sigma, \sigma_{int}, \sigma_{slope}$ = variance terms

Random slope and intercept example:

```
#Intercept varies with site, and slope of x can  
# also vary with site (both hierarchical)  
mm2 <- lmer(y ~ x + (x|site),data=dat)  
summary(mm2)
```

```
## Linear mixed model fit by REML ['lmerMod']  
## Formula: y ~ x + (x | site)  
## Data: dat  
##  
## REML criterion at convergence: 900.6  
##  
## Scaled residuals:  
##      Min       1Q   Median       3Q      Max   
## -2.10500 -0.64857  0.02414  0.61137  2.22996   
##  
## Random effects:  
## Groups   Name                Variance Std.Dev. Corr  
## site    (Intercept) 35.2210   5.9347  
##         x           0.7889   0.8882   0.82  
## Residual                9.3162   3.0522  
## Number of obs: 160, groups: site, 16  
##  
## Fixed effects:  
##              Estimate Std. Error t value  
## (Intercept)   2.0383    1.5091   1.351  
## x             0.6438    0.2275   2.830  
##  
## Correlation of Fixed Effects:  
##      (Intr)  
## x 0.790
```

Results from lmer model:

- Random effects:

Random slope and intercept example:

```
#Intercept varies with site, and slope of x can  
# also vary with site (both hierarchical)  
mm2 <- lmer(y ~ x + (x|site),data=dat)  
summary(mm2)
```

```
## Linear mixed model fit by REML ['lmerMod']  
## Formula: y ~ x + (x | site)  
## Data: dat  
##  
## REML criterion at convergence: 900.6  
##  
## Scaled residuals:  
##      Min       1Q   Median       3Q      Max   
## -2.10500 -0.64857  0.02414  0.61137  2.22996   
##  
## Random effects:  
## Groups   Name                Variance Std.Dev. Corr  
## site    (Intercept)  35.2210   5.9347  
##         x             0.7889   0.8882   0.82  
## Residual                9.3162   3.0522  
## Number of obs: 160, groups: site, 16  
##  
## Fixed effects:  
##              Estimate Std. Error t value  
## (Intercept)   2.0383    1.5091    1.351  
## x             0.6438    0.2275    2.830  
##  
## Correlation of Fixed Effects:  
##      (Intr)  
## x 0.790
```

Results from lmer model:

- Random effects:
 - *residual, slope, and site variance* (σ , σ_{int} , σ_{slope})

Random slope and intercept example:

```
#Intercept varies with site, and slope of x can  
# also vary with site (both hierarchical)  
mm2 <- lmer(y ~ x + (x|site),data=dat)  
summary(mm2)
```

```
## Linear mixed model fit by REML ['lmerMod']  
## Formula: y ~ x + (x | site)  
## Data: dat  
##  
## REML criterion at convergence: 900.6  
##  
## Scaled residuals:  
##      Min       1Q   Median       3Q      Max   
## -2.10500 -0.64857  0.02414  0.61137  2.22996   
##  
## Random effects:  
## Groups   Name                Variance Std.Dev. Corr  
## site    (Intercept)  35.2210   5.9347  
##         x           0.7889   0.8882   0.82  
## Residual                9.3162   3.0522  
## Number of obs: 160, groups: site, 16  
##  
## Fixed effects:  
##              Estimate Std. Error t value  
## (Intercept)   2.0383     1.5091   1.351  
## x             0.6438     0.2275   2.830  
##  
## Correlation of Fixed Effects:  
##      (Intr)  
## x 0.790
```

Results from lmer model:

- Random effects:
 - *residual, slope, and site variance* (σ , σ_{int} , σ_{slope})
 - Correlation b/w intercept and slope = 0.82

Random slope and intercept example:

```
#Intercept varies with site, and slope of x can  
# also vary with site (both hierarchical)  
mm2 <- lmer(y ~ x + (x|site),data=dat)  
summary(mm2)
```

```
## Linear mixed model fit by REML ['lmerMod']  
## Formula: y ~ x + (x | site)  
## Data: dat  
##  
## REML criterion at convergence: 900.6  
##  
## Scaled residuals:  
##      Min       1Q   Median       3Q      Max   
## -2.10500 -0.64857  0.02414  0.61137  2.22996   
##  
## Random effects:  
## Groups   Name                Variance Std.Dev. Corr  
## site    (Intercept) 35.2210   5.9347  
##         x           0.7889   0.8882   0.82  
## Residual                9.3162   3.0522  
## Number of obs: 160, groups: site, 16  
##  
## Fixed effects:  
##              Estimate Std. Error t value  
## (Intercept)   2.0383    1.5091    1.351  
## x             0.6438    0.2275    2.830  
##  
## Correlation of Fixed Effects:  
##      (Intr)  
## x 0.790
```

Results from lmer model:

- Random effects:
 - *residual, slope, and site* variance (σ , σ_{int} , σ_{slope})
 - Correlation b/w intercept and slope = 0.82
 - Sites with higher intercept *also* have a higher slope

Random slope and intercept example:

```
#Intercept varies with site, and slope of x can  
# also vary with site (both hierarchical)  
mm2 <- lmer(y ~ x + (x|site),data=dat)  
summary(mm2)
```

```
## Linear mixed model fit by REML ['lmerMod']  
## Formula: y ~ x + (x | site)  
## Data: dat  
##  
## REML criterion at convergence: 900.6  
##  
## Scaled residuals:  
##      Min       1Q   Median       3Q      Max   
## -2.10500 -0.64857  0.02414  0.61137  2.22996   
##  
## Random effects:  
## Groups   Name      Variance Std.Dev. Corr  
## site     (Intercept) 35.2210  5.9347  
##          x           0.7889  0.8882   0.82  
## Residual                9.3162  3.0522  
## Number of obs: 160, groups: site, 16  
##  
## Fixed effects:  
##              Estimate Std. Error t value  
## (Intercept)   2.0383    1.5091    1.351  
## x             0.6438    0.2275    2.830  
##  
## Correlation of Fixed Effects:  
##      (Intr)  
## x 0.790
```

Results from lmer model:

- Random effects:
 - *residual, slope, and site* variance (σ , σ_{int} , σ_{slope})
 - Correlation b/w intercept and slope = 0.82
 - Sites with higher intercept *also* have a higher slope
- Fixed effects:

Random slope and intercept example:

```
#Intercept varies with site, and slope of x can  
# also vary with site (both hierarchical)  
mm2 <- lmer(y ~ x + (x|site),data=dat)  
summary(mm2)
```

```
## Linear mixed model fit by REML ['lmerMod']  
## Formula: y ~ x + (x | site)  
## Data: dat  
##  
## REML criterion at convergence: 900.6  
##  
## Scaled residuals:  
##      Min       1Q   Median       3Q      Max   
## -2.10500 -0.64857  0.02414  0.61137  2.22996   
##  
## Random effects:  
## Groups   Name      Variance Std.Dev. Corr  
## site     (Intercept) 35.2210  5.9347  
##          x           0.7889  0.8882   0.82  
## Residual                9.3162  3.0522  
## Number of obs: 160, groups: site, 16  
##  
## Fixed effects:  
##              Estimate Std. Error t value  
## (Intercept)   2.0383    1.5091    1.351  
## x             0.6438    0.2275    2.830  
##  
## Correlation of Fixed Effects:  
##      (Intr)  
## x 0.790
```

Results from lmer model:

- Random effects:
 - *residual, slope, and site* variance (σ , σ_{int} , σ_{slope})
 - Correlation b/w intercept and slope = 0.82
 - Sites with higher intercept *also* have a higher slope
- Fixed effects:
 - Intercept and slope estimates

Model matrices

- X

```
##      (Intercept)      x
## 1              1 -4.248450
## 2              1  5.766103
## 3              1 -1.820462
## 4              1  7.660348
## 5              1  8.809346
## 6              1 -9.088870
```

```
##      sitea siteb sitec sited sitee sitef siteg siteh sitei sitej sitek sitel sitem
## 1         0      0      0      0      0      0      0      1      0      0      0      0      0
## 2         0      0      0      0      0      0      0      0      0      0      1      0      0
## 3         0      0      0      0      0      0      1      0      0      0      0      0      0
## 4         0      0      0      0      0      0      0      0      0      0      0      0      0
## 5         0      0      0      0      0      0      0      0      0      0      0      0      0
## 6         0      0      0      0      0      0      0      0      0      0      0      0      0
```

```
##      siten siteo sitep
## 1         0      0      0
## 2         0      0      0
## 3         0      0      0
## 4         1      0      0
## 5         0      1      0
## 6         0      0      1
```

```
##      sitea siteb sitec sited sitee      sitef      siteg siteh sitei      sitej sitek
## 1         0      0      0      0      0      0 0.000000  1.04215      0      0 0.000000      0
```

Model matrices

- X

```
##      (Intercept)      x
## 1              1 -4.248450
## 2              1  5.766103
## 3              1 -1.820462
## 4              1  7.660348
## 5              1  8.809346
## 6              1 -9.088870
```

- U

```
##      sitea siteb sitec sited sitee sitef siteg siteh sitei sitej sitek sitel sitem
## 1         0      0      0      0      0      0      0      1      0      0      0      0      0
## 2         0      0      0      0      0      0      0      0      0      0      1      0      0
## 3         0      0      0      0      0      0      1      0      0      0      0      0      0
## 4         0      0      0      0      0      0      0      0      0      0      0      0      0
## 5         0      0      0      0      0      0      0      0      0      0      0      0      0
## 6         0      0      0      0      0      0      0      0      0      0      0      0      0
##      siten siteo sitep
## 1         0      0      0
## 2         0      0      0
## 3         0      0      0
## 4         1      0      0
## 5         0      1      0
## 6         0      0      1
```

```
##      sitea siteb sitec sited sitee      sitef      siteg siteh sitei      sitej sitek
## 1         0      0      0      0      0      0 0.000000  1.04215      0      0 0.000000      0
```

Model matrices

- X

```
##      (Intercept)      x
## 1              1 -4.248450
## 2              1  5.766103
## 3              1 -1.820462
## 4              1  7.660348
## 5              1  8.809346
## 6              1 -9.088870
```

- U

```
##      sitea siteb sitec sited sitee sitef siteg siteh sitei sitej sitek sitel sitem
## 1         0      0      0      0      0      0      0      1      0      0      0      0      0
## 2         0      0      0      0      0      0      0      0      0      0      1      0      0
## 3         0      0      0      0      0      0      1      0      0      0      0      0      0
## 4         0      0      0      0      0      0      0      0      0      0      0      0      0
## 5         0      0      0      0      0      0      0      0      0      0      0      0      0
## 6         0      0      0      0      0      0      0      0      0      0      0      0      0
##      siten siteo sitep
## 1         0      0      0
## 2         0      0      0
## 3         0      0      0
## 4         1      0      0
## 5         0      1      0
## 6         0      0      1
```

- Ux

```
##      sitea siteb sitec sited sitee      sitef      siteg siteh sitei      sitej sitek
## 1         0      0      0      0      0      0 0.000000  1.04215      0      0 0.000000      0
```

Model matrices

- X

```
## (Intercept)      x
## 1           1 -4.248450
## 2           1  5.766103
## 3           1 -1.820462
## 4           1  7.660348
## 5           1  8.809346
## 6           1 -9.088870
```

- U


```
## sitea siteb sitec sited sitee sitef siteg siteh sitei sitej sitek sitel sitem
## 1     0     0     0     0     0     0     0     1     0     0     0     0     0
## 2     0     0     0     0     0     0     0     0     0     0     1     0     0
## 3     0     0     0     0     0     0     1     0     0     0     0     0     0
## 4     0     0     0     0     0     0     0     0     0     0     0     0     0
## 5     0     0     0     0     0     0     0     0     0     0     0     0     0
## 6     0     0     0     0     0     0     0     0     0     0     0     0     0
## siten siteo sitep
## 1     0     0     0
## 2     0     0     0
## 3     0     0     0
## 4     1     0     0
## 5     0     1     0
## 6     0     0     1
```

- Ux

```
## sitea siteb sitec sited sitee sitef siteg siteh sitei sitej sitek
## 1     0     0     0     0     0     0 0.000000 1.04215 0 0.000000 0
```


Mixed effect model results

- Regression line of x on y is allowed to move up or down around the main regression line for each site (random intercepts)




Mixed effect model results

- Regression line of x on y is allowed to move up or down around the main regression line for each site (random intercepts)
- Slope of regression line can be more or less steep for each site (random slopes)



Mixed effect model results

- Regression line of x on y is allowed to move up or down around the main regression line for each site (random intercepts)
- Slope of regression line can be more or less steep for each site (random slopes)
- Changes in intercepts and slopes are *normally distributed*, and in this example, are correlated with each other



Why do we need to do any of this?

“My supervisor told me to just use site as a fixed effect. Why can't I do that?”

- You can do it this way, but you may encounter the following problems:

Why do we need to do any of this?

“My supervisor told me to just use site as a fixed effect. Why can’t I do that?”

- You can do it this way, but you may encounter the following problems:
 - You lose the *partial pooling* that occurs in mixed effects models = Worse estimates of site effects!

Why do we need to do any of this?

“My supervisor told me to just use site as a fixed effect. Why can't I do that?”

- You can do it this way, but you may encounter the following problems:
 - You lose the *partial pooling* that occurs in mixed effects models = Worse estimates of site effects!
 - You lose 1 d.f. for each site = Type II error \uparrow = You may not find the fixed effect of interest, even if it's there!

Why do we need to do any of this?

“My supervisor told me to just use site as a fixed effect. Why can't I do that?”

- You can do it this way, but you may encounter the following problems:
 - You lose the *partial pooling* that occurs in mixed effects models = Worse estimates of site effects!
 - You lose 1 d.f. for each site = Type II error \uparrow = You may not find the fixed effect of interest, even if it's there!
 - Sites with low sample sizes may cause your models to break

Why do we need to do any of this?

“My supervisor told me to just use site as a fixed effect. Why can’t I do that?”

- You can do it this way, but you may encounter the following problems:
 - You lose the *partial pooling* that occurs in mixed effects models = Worse estimates of site effects!
 - You lose 1 d.f. for each site = Type II error \uparrow = You may not find the fixed effect of interest, even if it’s there!
 - Sites with low sample sizes may cause your models to break
 - People³ who have read statistics books published after 1980 may ask questions

³e.g. me

Why do we need to do any of this?

“My supervisor told me to just use site as a fixed effect. Why can’t I do that?”

- You can do it this way, but you may encounter the following problems:
 - You lose the *partial pooling* that occurs in mixed effects models = Worse estimates of site effects!
 - You lose 1 d.f. for each site = Type II error \uparrow = You may not find the fixed effect of interest, even if it’s there!
 - Sites with low sample sizes may cause your models to break
 - People³ who have read statistics books published after 1980 may ask questions
- However, if you have a low number of sites (1-10), fixed effects may work better

³e.g. me

Why do we need to do any of this?

“My supervisor told me to just use site as a fixed effect. Why can’t I do that?”

- You can do it this way, but you may encounter the following problems:
 - You lose the *partial pooling* that occurs in mixed effects models = Worse estimates of site effects!
 - You lose 1 d.f. for each site = Type II error \uparrow = You may not find the fixed effect of interest, even if it’s there!
 - Sites with low sample sizes may cause your models to break
 - People³ who have read statistics books published after 1980 may ask questions
- However, if you have a low number of sites (1-10), fixed effects may work better
 - Hard to estimate σ_{site} if number of sites is low

³e.g. me

Why do we need to do any of this?

“My supervisor told me to just use site as a fixed effect. Why can’t I do that?”

- You can do it this way, but you may encounter the following problems:
 - You lose the *partial pooling* that occurs in mixed effects models = Worse estimates of site effects!
 - You lose 1 d.f. for each site = Type II error \uparrow = You may not find the fixed effect of interest, even if it’s there!
 - Sites with low sample sizes may cause your models to break
 - People³ who have read statistics books published after 1980 may ask questions
- However, if you have a low number of sites (1-10), fixed effects may work better
 - Hard to estimate σ_{site} if number of sites is low
 - If stakes are high, it may be better to be more conservative about site intercepts

³e.g. me

Why do we need to do any of this?

“My supervisor told me to just use site as a fixed effect. Why can’t I do that?”

- You can do it this way, but you may encounter the following problems:
 - You lose the *partial pooling* that occurs in mixed effects models = Worse estimates of site effects!
 - You lose 1 d.f. for each site = Type II error \uparrow = You may not find the fixed effect of interest, even if it’s there!
 - Sites with low sample sizes may cause your models to break
 - People³ who have read statistics books published after 1980 may ask questions
- However, if you have a low number of sites (1-10), fixed effects may work better
 - Hard to estimate σ_{site} if number of sites is low
 - If stakes are high, it may be better to be more conservative about site intercepts
 - Easier to interpret (p-values, ANOVA, etc.)

³e.g. me

First challenge

- You're interested in how forest cover influences the mass of three bat species. Maybe some of the species do better in forests?

First challenge

- You're interested in how forest cover influences the mass of three bat species. Maybe some of the species do better in forests?
- You've weighed a bunch of bats across different forest covers (`batMass.csv`). However, these were collected across 15 separate sites. Perhaps some of the variation is just caused by the site?

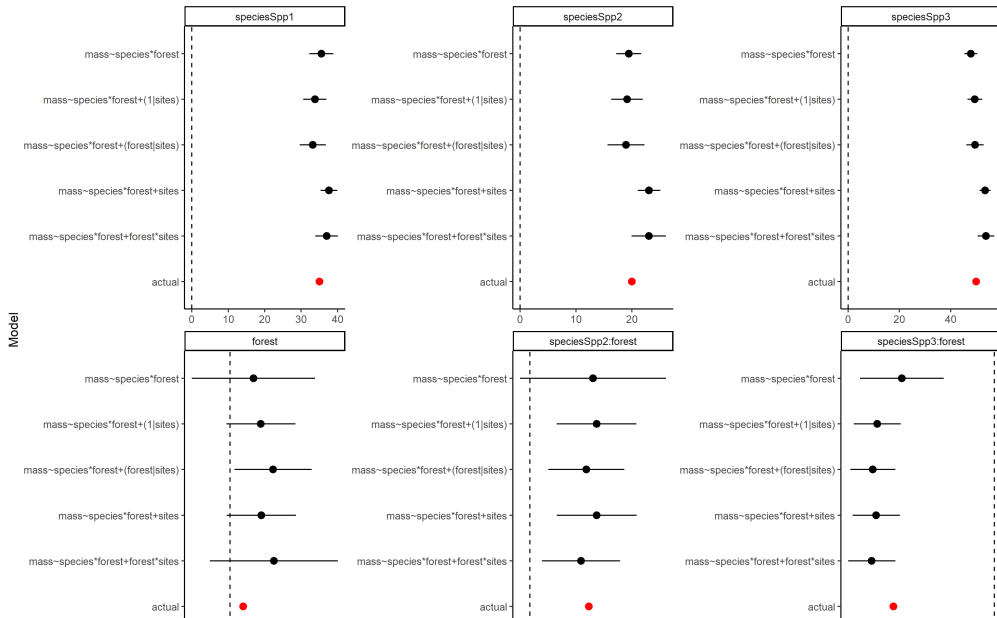
First challenge

- You're interested in how forest cover influences the mass of three bat species. Maybe some of the species do better in forests?
- You've weighed a bunch of bats across different forest covers (`batMass.csv`). However, these were collected across 15 separate sites. Perhaps some of the variation is just caused by the site?
- Fit a mixed effects model with the fixed effects you're interested in (**forest cover**, **species**), and include **site** as a random effect (*intercept* or *slope*)

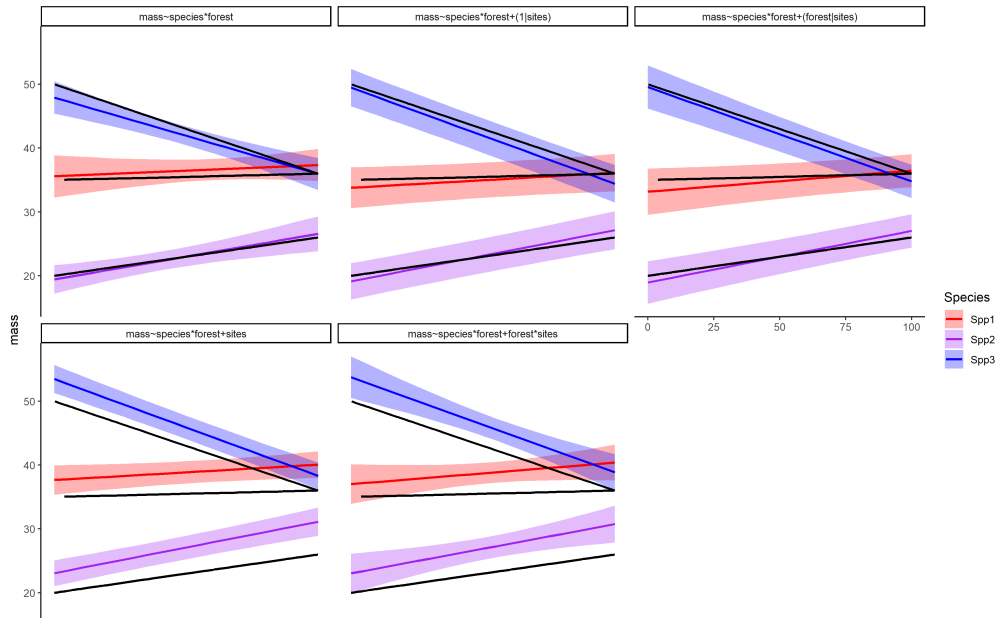
First challenge

- You're interested in how forest cover influences the mass of three bat species. Maybe some of the species do better in forests?
- You've weighed a bunch of bats across different forest covers (`batMass.csv`). However, these were collected across 15 separate sites. Perhaps some of the variation is just caused by the site?
- Fit a mixed effects model with the fixed effects you're interested in (**forest cover**, **species**), and include **site** as a random effect (*intercept* or *slope*)
- Your supervisor doesn't like hierarchical models, and tells you to just use site as another fixed term in an `lm` model. Do you get different results if you use their approach?

Results:



Results:



Part 2: GLMMs

What if my response variable is non-normal?

- Linear model (LM)

$$\hat{y} = X\beta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\hat{y} = X\beta + U\zeta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

$$\text{logit}(\hat{\phi}) = X\beta$$

$$y \sim \text{Binomial}(\hat{\phi})$$

$$\text{logit}(\hat{\phi}) = X\beta + U\zeta$$

$$y \sim \text{Binomial}(\hat{\phi})$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

What if my response variable is non-normal?

- Linear model (LM)

$$\hat{y} = X\beta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\hat{y} = X\beta + U\zeta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- Generalized linear model (GLM)

$$\text{logit}(\hat{\phi}) = X\beta$$

$$y \sim \text{Binomial}(\hat{\phi})$$

$$\text{logit}(\hat{\phi}) = X\beta + U\zeta$$

$$y \sim \text{Binomial}(\hat{\phi})$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

What if my response variable is non-normal?

- Linear model (LM)

$$\hat{y} = X\beta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- Linear mixed effects model (LMM)

$$\hat{y} = X\beta + U\zeta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- Generalized linear model (GLM)

$$\text{logit}(\hat{\phi}) = X\beta$$

$$y \sim \text{Binomial}(\hat{\phi})$$

$$\text{logit}(\hat{\phi}) = X\beta + U\zeta$$

$$y \sim \text{Binomial}(\hat{\phi})$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

What if my response variable is non-normal?

- Linear model (LM)

$$\hat{y} = X\beta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

- Linear mixed effects model (LMM)

$$\hat{y} = X\beta + U\zeta$$

$$y \sim \text{Normal}(\hat{y}, \sigma)$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

- Generalized linear model (GLM)

$$\text{logit}(\hat{\phi}) = X\beta$$

$$y \sim \text{Binomial}(\hat{\phi})$$

- Generalized linear mixed effects model (GLMM)

$$\text{logit}(\hat{\phi}) = X\beta + U\zeta$$

$$y \sim \text{Binomial}(\hat{\phi})$$

$$\zeta \sim \text{Normal}(0, \sigma_{\text{site}})$$

How do I fit GLMMs?

- `glmer` and `glmer.nb` from `library(lme4)` work for Binomial, Poisson, and Negative Binomial data

```
library(lme4)
glmm1 <- glmer.nb(y~x+(x|site),data=dat2) #Negative binomial GLMM
summary(glmm1) #glmer.nb takes a long time to run
```

```
library(glmmTMB)
glmm2 <- glmmTMB(y~x+(x|site),data=dat2,family=nbinom2())
summary(glmm2) #Similar results, but quicker
```

```
## Family: nbinom2 ( log )
## Formula:          y ~ x + (x | site)
## Data: dat2
##
##      AIC      BIC   logLik deviance df.resid
##    627.8    646.2   -307.9    615.8     154
##
## Random effects:
##
## Conditional model:
##   Groups Name      Variance Std.Dev. Corr
## site  (Intercept) 1.43543  1.1981
##      x             0.02892  0.1701  0.92
## ---
## Null model deviance: 160
## Full model deviance: 154
```

How do I fit GLMMs?

- `glmer` and `glmer.nb` from `library(lme4)` work for Binomial, Poisson, and Negative Binomial data

```
library(lme4)
glmm1 <- glmer.nb(y~x+(x|site),data=dat2) #Negative binomial GLMM
summary(glmm1) #glmer.nb takes a long time to run
```

- `glmmTMB` from `library(glmmTMB)` works for those above, *plus* a bunch of others

```
library(glmmTMB)
glmm2 <- glmmTMB(y~x+(x|site),data=dat2,family=nbinom2())
summary(glmm2) #Similar results, but quicker
```

```
## Family: nbinom2 ( log )
## Formula:          y ~ x + (x | site)
## Data: dat2
##
##      AIC      BIC   logLik deviance df.resid
##    627.8    646.2   -307.9    615.8      154
##
## Random effects:
##
## Conditional model:
##   Groups Name      Variance Std.Dev. Corr
## site  (Intercept) 1.43543   1.1981
##      x            0.02892   0.1701  0.92
## Number of parameters: 16
```

How do I fit GLMMs?

- `glmer` and `glmer.nb` from `library(lme4)` work for Binomial, Poisson, and Negative Binomial data

```
library(lme4)
glmm1 <- glmer.nb(y~x+(x|site),data=dat2) #Negative binomial GLMM
summary(glmm1) #glmer.nb takes a long time to run
```

- `glmmTMB` from `library(glmmTMB)` works for those above, *plus* a bunch of others
 - e.g. Zero-inflation, Beta-binomial, Spatial Models

```
library(glmmTMB)
glmm2 <- glmmTMB(y~x+(x|site),data=dat2,family=nbinom2())
summary(glmm2) #Similar results, but quicker
```

```
## Family: nbinom2 ( log )
## Formula: y ~ x + (x | site)
## Data: dat2
##
##      AIC      BIC   logLik deviance df.resid
##    627.8    646.2   -307.9    615.8      154
##
## Random effects:
##
## Conditional model:
##   Groups Name      Variance Std.Dev. Corr
##   site  (Intercept) 1.43543  1.1981
##         x           0.02892  0.1701  0.92
##
```

Fitting GLMMs - glmer.nb

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
##   Approximation) [glmerMod]
##   Family: Negative Binomial(5.1294)  ( log )
## Formula: y ~ x + (x | site)
##   Data: dat2
##
##           AIC      BIC   logLik deviance df.resid
##      627.8    646.3   -307.9   615.8      154
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.3745 -0.7098 -0.3946  0.5108  2.5367
##
## Random effects:
##   Groups Name      Variance Std.Dev. Corr
##   site  (Intercept) 1.43500  1.1979
##         x           0.02878  0.1697   0.92
## Number of obs: 160, groups:  site, 16
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.32746    0.32166   1.018  0.3087
## x            0.10830    0.04681   2.314  0.0207 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##   (Intr)
## x 0.799
```

Fitting GLMMs - glmmTMB


```
## Family: nbinom2 ( log )
## Formula:          y ~ x + (x | site)
## Data: dat2
##
##      AIC      BIC   logLik deviance df.resid
##    627.8    646.2   -307.9   615.8     154
##
## Random effects:
##
## Conditional model:
##   Groups Name      Variance Std.Dev. Corr
##   site  (Intercept) 1.43543  1.1981
##         x           0.02892  0.1701  0.92
## Number of obs: 160, groups: site, 16
##
## Dispersion parameter for nbinom2 family (): 5.12
##
## Conditional model:
##           Estimate Std. Error z value Pr(>|z|)
## (Intercept) 0.34132    0.32172   1.061  0.2887
## x           0.11026    0.04697   2.348  0.0189 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


Residual checks on glmmTMB models


- Extract residuals and make your own plots, or use `simulateResiduals` from `library(DHARMA)` (see [here](#))

DHARMA residual

QQ plot residuals




Residual vs. predicted
No significant problems detected



Residual checks on glmmTMB models


- Extract residuals and make your own plots, or use `simulateResiduals` from `library(DHARMA)` (see [here](#))
- DHARMA also has useful functions for checking overdispersion and zero-inflation (found [here](#))

DHARMA residual



Partial residual plots for glmmTMB models

- `ggpredict()` from `library(ggeffects)` works with `glmmTMB` models



Part 3: Hypothesis testing and inference

I fit a model... now what?

- Congratulations, your model ran and it met the assumptions of regression...

I fit a model... now what?

- Congratulations, your model ran and it met the assumptions of regression...
- Time to see if your predictions are supported by your data!

I fit a model... now what?

- Congratulations, your model ran and it met the assumptions of regression...
- Time to see if your predictions are supported by your data!
- For each of the terms in your model:

I fit a model... now what?

- Congratulations, your model ran and it met the assumptions of regression...
- Time to see if your predictions are supported by your data!
- For each of the terms in your model:
 - Was the term “important” in your model?

I fit a model... now what?

- Congratulations, your model ran and it met the assumptions of regression...
- Time to see if your predictions are supported by your data!
- For each of the terms in your model:
 - Was the term “important” in your model?
 - If so, what direction was the effect in?

I fit a model... now what?

- Congratulations, your model ran and it met the assumptions of regression...
- Time to see if your predictions are supported by your data!
- For each of the terms in your model:
 - Was the term “important” in your model?
 - If so, what direction was the effect in?
- How well did your model fit your data (overall)?

I fit a model... now what?

- Congratulations, your model ran and it met the assumptions of regression...
- Time to see if your predictions are supported by your data!
- For each of the terms in your model:
 - Was the term “important” in your model?
 - If so, what direction was the effect in?
- How well did your model fit your data (overall)?
- Some other bits of advice

Step 1: was the term “important”?

In linear models, this is done using an ANOVA F-test (also shown at the bottom of a `summary()` statement):

```
## Analysis of Variance Table
##
## Response: mpg
##      Df Sum Sq Mean Sq F value    Pr(>F)
## disp    1 808.89   808.89 73.9959 1.788e-09 ***
## gear    1   0.14    0.14  0.0132   0.9093
## Residuals 29 317.01   10.93
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

But wait, there's more!

Unfortunately, if we have more than 1 term in the model, the order of terms can change your answer:

- ```
Analysis of Variance Table
##
Response: mpg
Df Sum Sq Mean Sq F value Pr(>F)
disp 1 808.89 808.89 73.9959 1.788e-09 ***
gear 1 0.14 0.14 0.0132 0.9093
Residuals 29 317.01 10.93

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## But wait, there's more!

Unfortunately, if we have more than 1 term in the model, the order of terms can change your answer:

- ```
## Analysis of Variance Table
##
## Response: mpg
##           Df Sum Sq Mean Sq F value    Pr(>F)
## disp       1  808.89   808.89  73.9959 1.788e-09 ***
## gear       1    0.14     0.14   0.0132   0.9093
## Residuals 29  317.01    10.93
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```
- ```
Analysis of Variance Table
##
Response: mpg
Df Sum Sq Mean Sq F value Pr(>F)
gear 1 259.75 259.75 23.762 3.595e-05 ***
disp 1 549.28 549.28 50.248 8.465e-08 ***
Residuals 29 317.01 10.93

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## Solution: use a Type II ANOVA

- We're usually interested in the importance of a term on its own, *not just after other terms are accounted for* (Type I ANOVA).

## Solution: use a Type II ANOVA

- We're usually interested in the importance of a term on its own, *not just after other terms are accounted for* (Type I ANOVA).
- For this, we use a Type II ANOVA. This can be done using `drop1()` (or `Anova()` in the `car` package)



## Solution: use a Type II ANOVA

- We're usually interested in the importance of a term on its own, *not just after other terms are accounted for* (Type I ANOVA).
- For this, we use a Type II ANOVA. This can be done using `drop1()` (or `Anova()` in the `car` package)

- ```
## Single term deletions
##
## Model:
## mpg ~ disp + gear
##      Df Sum of Sq    RSS      AIC F value    Pr(>F)
## <none>                 317.01   79.383
## disp    1    549.28 866.30 109.552 50.2476 8.465e-08 ***
## gear    1      0.14 317.16  77.397  0.0132  0.9093
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Solution: use a Type II ANOVA

- We're usually interested in the importance of a term on its own, *not just after other terms are accounted for* (Type I ANOVA).
- For this, we use a Type II ANOVA. This can be done using `drop1()` (or `Anova()` in the `car` package)

- ```
Single term deletions
##
Model:
mpg ~ disp + gear
Df Sum of Sq RSS AIC F value Pr(>F)
<none> 317.01 79.383
disp 1 549.28 866.30 109.552 50.2476 8.465e-08 ***
gear 1 0.14 317.16 77.397 0.0132 0.9093

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- ```
## Single term deletions
##
## Model:
## mpg ~ gear + disp
##      Df Sum of Sq    RSS      AIC F value    Pr(>F)
## <none>                 317.01   79.383
## gear    1      0.14 317.16  77.397  0.0132  0.9093
## disp    1    549.28 866.30 109.552 50.2476 8.465e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Interactions make ANOVA testing a bit strange

- If interactions are present, it doesn't really make sense to test the main terms **because they depend on the interactions**

Interactions make ANOVA testing a bit strange

- If interactions are present, it doesn't really make sense to test the main terms **because they depend on the interactions**
- For example:

```
## Single term deletions
##
## Model:
## mpg ~ gear * disp
##           Df Sum of Sq    RSS    AIC F value    Pr(>F)
## <none>                 278.60  77.250
## gear:disp  1     38.412  317.01  79.383   3.8604 0.05943 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Interactions make ANOVA testing a bit strange

- If interactions are present, it doesn't really make sense to test the main terms **because they depend on the interactions**
- For example:

```
## Single term deletions
##
## Model:
## mpg ~ gear * disp
##      Df Sum of Sq    RSS    AIC F value    Pr(>F)
## <none>                 278.60  77.250
## gear:disp  1      38.412 317.01  79.383   3.8604 0.05943 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Anova from the car package tests other terms *after* interactions, but the meaning isn't the same. I prefer to keep things simple and just use drop1()

```
## Anova Table (Type II tests)
##
## Response: mpg
##      Sum Sq Df F value    Pr(>F)
## gear      0.14  1  0.0145   0.90502
## disp    549.28  1 55.2038 4.312e-08 ***
## gear:disp  38.41  1   3.8604  0.05943 .
## Residuals 278.60 28
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

GLMs, LMMs, and GLMMs

- drop1 also works with GLMs, LMMs, and GLMMs, but we use a χ^2 *likelihood ratio test* rather than an F-test

GLMs, LMMs, and GLMMs

- drop1 also works with GLMs, LMMs, and GLMMs, but we use a χ^2 *likelihood ratio test* rather than an F-test
- Unfortunately, different numbers of data points change your likelihood, so this can wreck LRTs. *This happens a lot if you have NAs in your predictor columns*, so clean up your data before using it in models.

ML vs REML

- Maximum likelihood (ML) estimates of variance (e.g. SD) are always smaller than the actual variance (biased)

Solution:

ML vs REML

- Maximum likelihood (ML) estimates of variance (e.g. SD) are always smaller than the actual variance (biased)
- Restricted maximum likelihood (REML) uses a mathematical trick to get around this, but...

Solution:

ML vs REML

- Maximum likelihood (ML) estimates of variance (e.g. SD) are always smaller than the actual variance (biased)
- Restricted maximum likelihood (REML) uses a mathematical trick to get around this, but. . .
- This means that models with different numbers of terms don't have the same REML estimates

Solution:

ML vs REML

- Maximum likelihood (ML) estimates of variance (e.g. SD) are always smaller than the actual variance (biased)
- Restricted maximum likelihood (REML) uses a mathematical trick to get around this, but...
- This means that models with different numbers of terms don't have the same REML estimates
- Likelihood between these models technically can't be compared!

Solution:

ML vs REML

- Maximum likelihood (ML) estimates of variance (e.g. SD) are always smaller than the actual variance (biased)
- Restricted maximum likelihood (REML) uses a mathematical trick to get around this, but...
- This means that models with different numbers of terms don't have the same REML estimates
- Likelihood between these models technically can't be compared!

Solution:

- ① Use ML if comparing between models with different fixed effects, then...

ML vs REML

- Maximum likelihood (ML) estimates of variance (e.g. SD) are always smaller than the actual variance (biased)
- Restricted maximum likelihood (REML) uses a mathematical trick to get around this, but...
- This means that models with different numbers of terms don't have the same REML estimates
- Likelihood between these models technically can't be compared!

Solution:

- ① Use ML if comparing between models with different fixed effects, then...
- ② Re-fit with REML once you've decided on a model

Model selection

How many terms should be in my model?

- Same principle as in regular linear models: **what do you think the process is?**

⁴"My God, it's full of stars!" -2001, A Space Odyssey

Model selection

How many terms should be in my model?

- Same principle as in regular linear models: **what do you think the process is?**
 - Just because a term is “not significant” doesn't mean it should be dropped out!

⁴"My God, it's full of stars!" -2001, A Space Odyssey

Model selection

How many terms should be in my model?

- Same principle as in regular linear models: **what do you think the process is?**
 - Just because a term is “not significant” doesn’t mean it should be dropped out!
 - Just because a term is “significant” doesn’t mean it should be left in!

⁴“My God, it’s full of stars!” -2001, A Space Odyssey

Model selection

How many terms should be in my model?

- Same principle as in regular linear models: **what do you think the process is?**
 - Just because a term is “not significant” doesn’t mean it should be dropped out!
 - Just because a term is “significant” doesn’t mean it should be left in!
 - Avoid selecting models based on R^2 or AIC alone. Avoid stargazing⁴(hunting around for “better” p-values or AIC scores)

⁴"My God, it's full of stars!" -2001, A Space Odyssey

Model selection

How many terms should be in my model?

- Same principle as in regular linear models: **what do you think the process is?**
 - Just because a term is “not significant” doesn’t mean it should be dropped out!
 - Just because a term is “significant” doesn’t mean it should be left in!
 - Avoid selecting models based on R^2 or AIC alone. Avoid stargazing⁴(hunting around for “better” p-values or AIC scores)
- To test whether terms are important in predicting your data (similar to), use *likelihood-ratio tests*

⁴"My God, it's full of stars!" -2001, A Space Odyssey

Model selection

How many terms should be in my model?

- Same principle as in regular linear models: **what do you think the process is?**
 - Just because a term is “not significant” doesn’t mean it should be dropped out!
 - Just because a term is “significant” doesn’t mean it should be left in!
 - Avoid selecting models based on R^2 or AIC alone. Avoid stargazing⁴(hunting around for “better” p-values or AIC scores)
- To test whether terms are important in predicting your data (similar to), use *likelihood-ratio tests*
 - `drop1(model, test='Chisq')`

⁴"My God, it's full of stars!" -2001, A Space Odyssey

Model selection

How many terms should be in my model?

- Same principle as in regular linear models: **what do you think the process is?**
 - Just because a term is “not significant” doesn’t mean it should be dropped out!
 - Just because a term is “significant” doesn’t mean it should be left in!
 - Avoid selecting models based on R^2 or AIC alone. Avoid stargazing⁴(hunting around for “better” p-values or AIC scores)
- To test whether terms are important in predicting your data (similar to), use *likelihood-ratio tests*
 - `drop1(model, test='Chisq')`
 - AIC tests usually say the same thing as LR tests

⁴"My God, it's full of stars!" -2001, A Space Odyssey